

Secure Programming, Continued
Computer Security
Peter Reiher
May 17, 2016

Outline

- Introduction
- Principles for secure software
- **Major problem areas**

Example Problem Areas

- Buffer overflows
- Error handling
- Privilege escalation
- Race conditions
- Use of randomness
- Proper use of cryptography
- Trust
- Input verification
- Variable synchronization
- Variable initialization

Error Handling

- Error handling code often gives attackers great possibilities
- It's rarely executed and often untested
- So it might have undetected errors
- Attackers often try to compromise systems by forcing errors

A Typical Error Handling Problem

- Not cleaning everything up
- On error conditions, some variables don't get reset
- If error not totally fatal, program continues with old values
- Could cause security mistakes
 - E.g., not releasing privileges when you should

Some Examples

- Remote denial of service attack on Apache HTTP server due to bad error handling (2010)
- ntpd (Network Time Protocol Daemon) error handling flaw (2015)
 - Essentially allowing attacker to set target's clock

Checking Return Codes

- A generalization of error handling
- Always check return codes
- A security program manager for Microsoft said this is his biggest problem
- Very dangerous to bull ahead if it turns out your call didn't work properly
- Example: Nagios XI didn't check the return value of `setuid()` call, allowing privilege escalation

Privilege Escalation

- Programs usually run under their user's identity with his privileges
- Some programs get expanded privileges
 - Setuid programs in Unix, e.g.
- Poor programming here can give too much access

An Example Problem

- A program that runs setuid and allows a shell to be forked
 - Giving the caller a root environment in which to run arbitrary commands
- Buffer overflows in privileged programs usually give privileged access
- *Privilege escalation* – using program flaws to obtain greater access privileges you shouldn't get

A Real World Example

- **Lenovo System Update Service**
- Run by an unprivileged user to perform valid system updates
- Created a temporary admin account that had privileges to perform these updates
- Account name and password were predictable, allowing attacker to guess them
- And run as system administrator

What To Do About This?

- Avoid running programs setuid
 - Or in other OSs' high privilege modes
- If you must, don't make them root-owned
 - Remember, least privilege
- Change back to the real caller as soon as you can
 - Limiting exposure
- Use virtualization to compartmentalize

Virtualization Approaches

- Run stuff in a virtual machine
 - Only giving access to safe stuff
- Hard to specify what's safe
- Hard to allow safe interactions between different VMs
- VM might not have perfect isolation

Race Conditions

- A common cause of security bugs
- Usually involve multiprocessing or multithreaded programs
- Caused by different threads of control operating in unpredictable fashion
 - When programmer thought they'd work in a particular order

What Is a Race Condition?

- A situation in which two (or more) threads of control are cooperating or sharing something
- If their events happen in one order, one thing happens
- If their events happen in another order, something else happens
- Often the results are unforeseen

Security Implications of Race Conditions

- Usually you checked privileges at one point
- You thought the next lines of code would run next
 - So privileges still apply
- But multiprocessing allows things to happen in between

The TOCTOU Issue

- Time of Check to Time of Use
- Have security conditions changed between when you checked?
- And when you used it?
- Multiprogramming issues can make that happen
- Sometimes under attacker control

A Short Detour

- In Unix, processes can have two associated user IDs
 - Effective ID
 - Real ID
- Real ID is the ID of the user who actually ran it
- Effective ID is current ID for access control purposes
- Setuid programs run this way
- System calls allow you to manipulate it

Effective UID and Access Permissions

- Unix checks accesses against effective UID, not real UID
- So setuid program uses permissions for the program's owner
 - Unless relinquished
- Remember, root has universal access privileges

An Example

- Code from Unix involving a temporary file
- Runs setuid root

```
res = access("/tmp/userfile", R_OK);  
If (res != 0)  
    die("access");  
fd = open("/tmp/userfile", O_RDONLY);
```

What's (Supposed to Be) Going on Here?

- Checked access on `/tmp/userfile` to make sure user was allowed to read it
 - User can use links to control what this file is
- `access()` checks real user ID, not effective one
 - So checks access permissions not as root, but as actual user
- So if user can read it, open file for read
 - Which root is definitely allowed to do
- Otherwise exit

What's Really Going On Here?

- This program might not run uninterrupted
- OS might schedule something else in the middle
- In particular, between those two lines of code

How the Attack Works

- Attacker puts innocuous file in
 /tmp/userfile
- Calls the program
- Quickly deletes file and replaces it with
link to sensitive file
 - One only readable by root
- If timing works, he gets secret contents



The Dynamics of the Attack

Success!

Let's try that again!

/etc/secretfile

/tmp/userfile



1. Run program
2. Change file

```
res = access("/tmp/userfile", R_OK);  
if (res != 0)  
    die("access");  
fd = open("/tmp/userfile", O_RDONLY);
```



How Likely Was That?

- Not very
 - The timing had to be just right
- But the attacker can try it many times
 - And may be able to influence system to make it more likely
- And he only needs to get it right once
- Timing attacks of this kind can work
- The longer between check and use, the more dangerous

Some Types of Race Conditions

- **File races**
 - Which file you access gets changed
- **Permissions races**
 - File permissions are changed
- **Ownership races**
 - Who owns a file changes
- **Directory races**
 - Directory hierarchy structure changes

A Real Example

- In the Linux SystemTap utility (2011)
 - Which gathers info about the system
- Allows modules to be loaded
- Checks privileges, but then there's a delay before loading the module
 - So it might load a different module
 - Allowing privilege escalation
- A genuine TOCTOU issue

Preventing Race Conditions

- Minimize time between security checks and when action is taken
- Be especially careful with files that users can change
- Use locking and features that prevent interruption, when possible
- Avoid designs that require actions where races can occur

Randomness and Determinism

- Many pieces of code require some randomness in behavior
- Where do they get it?
- As earlier key generation discussion showed, it's not that easy to get

Pseudorandom Number Generators

- PRNG
- Mathematical methods designed to produce strings of random-like numbers
- Actually deterministic
 - But share many properties with true random streams of numbers

Attacks on PRNGs

- Cryptographic attacks
 - Observe stream of numbers and try to deduce the function
- State attacks
 - Attackers gain knowledge of or influence the internal state of the PRNG

An Example

- ASF Software's Texas Hold'Em Poker
- Flaw in PRNG allowed cheater to determine everyone's cards
 - Flaw in card shuffling algorithm
 - Seeded with a clock value that can be easily obtained

Another Example

- Netscape's early SSL implementation
- Another guessable seed problem
 - Based on knowing time of day, process ID, and parent process ID
 - Process IDs readily available by other processes on same box
- Broke keys in 30 seconds

A Recent Case

- The chip-and-pin system is used to secure smart ATM cards
- Uses cryptographic techniques that require pseudo-random numbers
- Cambridge found weaknesses in the PRNG
- Allows attackers to withdraw cash without your card
- Seems to be in real use in the wild

How to Do Better?

- Use hardware randomness, where available
- Use high quality PRNGs
 - Preferably based on entropy collection methods
- Don't use seed values obtainable outside the program

Proper Use of Cryptography

- Never write your own crypto functions if you have any choice
 - Another favorite piece of advice from industry
- Never, ever, design your own encryption algorithm
 - Unless that's your area of expertise
- Generally, rely on tried and true stuff
 - Both algorithms and implementations

Proper Use of Crypto

- Even with good crypto algorithms (and code), problems are possible
- Proper use of crypto is quite subtle
- Bugs possible in:
 - Choice of keys
 - Key management
 - Application of cryptographic ops

An Example

- An application where RSA was used to distribute a triple-DES key
- Seemed to work fine
- Someone noticed that part of the RSA key exchange was always the same
 - That's odd . . .

What Was Happening?

- Bad parameters were handed to the RSA encryption code
- It failed and returned an error
- Which wasn't checked for
 - Since it “couldn't fail”
- As a result, RSA encryption wasn't applied at all
- The session key was sent in plaintext . . .

Another Example

- From an Android app
- It derived an encryption key from the user's password
 - The key looked like this:
 - EFBFBDEFBFBDEFBFBFD603466EFBFBFD7BEFBFBFD6C24E2B2AA576AEFBFBDEFBFBDEFBFBFD0C6BEFBFBDEFBFBDEFBFBDEFBFBFD76EFBFBDEFBFBDEFBFBDEFBFBDEFBFBDEFBFBFD
 - Hmm . . .
- They created the key as a byte array
 - So far, so good
- Then cast the byte array to a string

Why Did That Cause Problems?

- Android's default character set is UTF-8
- If UTF-8 gets a bit pattern that isn't a proper character, it replaces it with
 - The hex string “EFBFBD”
- Lots of random bit patterns aren't UTF characters, so

B7B0F88D603466CF7BF26C24E2B2AA576AAFC5E90C6BD4EECCC576B9D7F1E9C3

Was cast to

EFBFBDDEFBFBDEFBFBBD603466EFBFBBD7BEFBFBBD6C24E2B2AA576AEFBFBDEFBFBDEFBFBBD0C6B
EFBFBDDEFBFBDEFBFBDEFBFBBD76EFBFBDEFBFBDEFBFBDEFBFBDEFBFBBD

Trust Management

- Don't trust anything you don't need to
- Don't trust other programs
- Don't trust other components of your program
- Don't trust users
- Don't trust the data users provide you

Trust

- Some trust required to get most jobs done
- But determine how much you must trust the other
 - Don't trust things you can independently verify
- Limit the scope of your trust
 - Compartmentalization helps
- Be careful who you trust

Two Important Lessons

1. Many security problems arise because of unverified assumptions
 - You think someone is going to do something he actually isn't
2. Trusting someone doesn't just mean trusting their honesty
 - It means trusting their caution, too

Input Verification

- Never assume users followed any rules in providing you input
- They can provide you with anything
- Unless you check it, assume they've given you garbage
 - Or worse
- Just because the last input was good doesn't mean the next one will be

Treat Input as Hostile

- If it comes from outside your control and reasonable area of trust
- Probably even if it doesn't
- There may be code paths you haven't considered
- New code paths might be added
- Input might come from new sources

For Example

- Shopping cart exploits
- Web shopping carts sometimes handled as a cookie delivered to the user
- Some of these weren't encrypted
- So users could alter them
- The shopping cart cookie included the price of the goods . . .

What Was the Problem?

- The system trusted the shopping cart cookie when it was returned
 - When there was no reason to trust it
- Either encrypt the cookie
 - Making the input more trusted
 - Can you see any problem with this approach?
- Or scan the input before taking action on it
 - To find refrigerators being sold for 3 cents

Variable Synchronization

- Often, two or more program variables have related values
- Common example is a pointer to a buffer and a length variable
- Are the two variables always synchronized?
- If not, bad input can cause trouble

An Example

- From Apache web server
- `cdata` is a pointer to a buffer
- `len` is an integer containing the length of that buffer
- Programmer wanted to get rid of leading and trailing white spaces

The Problematic Code

```
while (apr_isspace(*cdata))  
    ++cdata;  
while (len-- >0 &&  
    apr_isspace(cdata[len]))  
    continue;  
cdata[len+1] = '\0';
```

- len is not decremented when leading white spaces are removed
- So trailing white space removal can overwrite end of buffer with nulls
- May or may not be serious security problem, depending on what's stored in overwritten area

Variable Initialization

- Some languages let you declare variables without specifying their initial values
- And let you use them without initializing them
 - E.g., C and C++
- Why is that a problem?

Variable Initialization

- Some languages let you declare variables without specifying their initial values
- And let you use them without initializing them
 - E.g., C and C++
- Why is that a problem?

A Little Example

```
main()
{
    foo();
    bar();
}

foo()
{
    int a;
    int b;
    int c;

    a = 11;
    b = 12;
    c = 13;
}

bar()
{
    int aa;
    int bb;
    int cc;

    printf("aa = %d\n", aa);
    printf("bb = %d\n", bb);
    printf("cc = %d\n", cc);
}
```

What's the Output?

```
lever.cs.ucla.edu[9] ./a.out
```

```
aa = 11
```

```
bb = 12
```

```
cc = 13
```

- Perhaps not exactly what you might want

Why Is This Dangerous?

- Values from one function “leak” into another function
- If attacker can influence the values in the first function,
- Maybe he can alter the behavior of the second one

Variable Cleanup

- Often, programs reuse a buffer or other memory area
- If old data lives in this area, might not be properly cleaned up
- And then can be treated as something other than what it really was
- E.g., bug in Microsoft TCP/IP stack
 - Old packet data treated as a function pointer

Use-After-Free Bugs

- Increasingly popular security bug type
- Related to memory management
 - Memory structures are dynamically allocated on the heap
- Either explicitly or implicitly freed
 - Depending on language and context
- In some cases, pointers can be used to access freed memory
 - E.g., in C and C++

An Example Use-After-Free Bug

- In OpenSSL (from 2009)

```
...
frag->fragment, frag->msg_header.frag_len);
}
dtls1_tm_fragment_free(frag);
dtls1_tm_fragment_free(frag);
pitem_free(item);

if (al==0)
{
    *ok = 1;
    return frag->msg_header.frag_len;
    return frag->msg_header.frag_len
}
```

What Was the Effect?

- Typically, crashing the program
- But it would depend
- When combined with other vulnerabilities, could be worse
- E.g., arbitrary code execution
- Often making use of poor error handling code

Recent Examples of Use-After-Free Bugs

- Internet Explorer (2014, several in 2012-2013)
- Adobe Flash (2016, multiple cases in 2015)
- Mozilla, multiple products (2012)
- Google Chrome (2012)

Some Other Problem Areas

- Handling of data structures
 - Indexing error in DAEMON Tools
- Arithmetic issues
 - Integer overflow in Adobe Flash (2016)
 - Signedness error in XnView (2012)
- Errors in flow control
 - Samba error that causes loop to use wrong structure
- Off-by-one errors
 - Denial of service flaw in Clam AV (2011)

Yet More Problem Areas

- Null pointer dereferencing
 - FreeBSD denial of service (2016)
- Side effects
- Punctuation errors
- Typos and cut-and-paste errors
 - iOS vulnerability based on inadvertent duplication of a goto statement (2014)
- There are many others

Why Should You Care?

- A lot of this stuff is kind of exotic
- Might seem unlikely it can be exploited
- Sounds like it would be hard to exploit without source code access
- Many examples of these bugs probably unexploitable

So . . . ?

- Well, that's what everyone thinks before they get screwed
- “Nobody will find this bug”
- “It's too hard to figure out how to exploit this bug”
- “It will get taken care of by someone else”
 - Code auditors
 - Testers
 - Firewalls

That's What They Always Say

- Before their system gets screwed
- Attackers can be very clever
 - Maybe more clever than you
- Attackers can work very hard
 - Maybe harder than you would
- Attackers may not have the goals you predict

But How to Balance Things?

- You only have a certain amount of time to design and build code
- Won't secure coding cut into that time?
- Maybe
- But less if you develop code coding practices
- If you avoid problematic things, you'll tend to code more securely

Some Good Coding Practices

- Validate input
- Be careful with failure conditions and return codes
- Avoid dangerous constructs
 - Like C input functions that don't specify amount of data
- Keep it simple