

Layers, Naming, and Sockets

CS 118

Computer Network Fundamentals

Peter Reiher


Outline

- What's a party?
- Inside names
- Outside names
- Linking the two
- Sockets as an example

Recall: definitions

- Communication
 - Methods for exchanging information between a fixed set of directly-connected parties using a single protocol
- Networking
 - Methods to enable communication between varying sets of indirectly connected parties that don't share a single protocol
- Protocol
 - A set of rules, agreed in advance [between the parties], that enable communication

Names, layers, and translations




Today's theme: #WTPA?

- What is a “party”?
- Where is that party?
 - Physically (in some physical place)
 - Logically (in some layer)


A network layer

- Nodes
 - Sources and sinks of information
- Links
 - Channels that connect two or more nodes



A closer look:

- What's inside a node?
 - What actually communicates to the outside?



One View




Google




One node has channels
to multiple parties

Another view

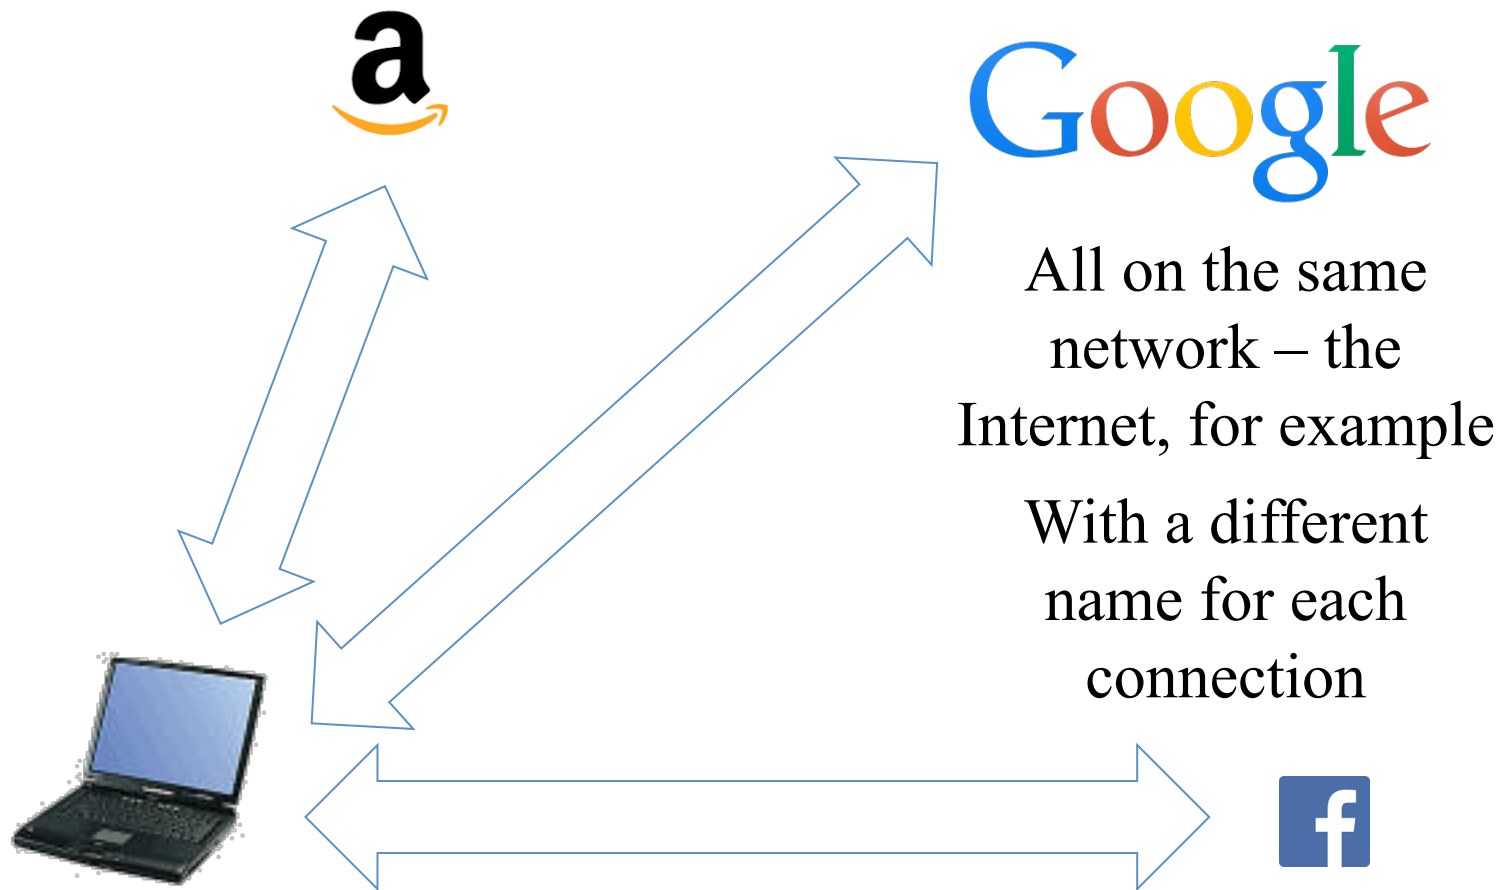


A closer look v2:

- What's inside a node when:
 - It has multiple channels on a single network (several names used external to the node)?




What Does That Mean?




A closer look v3:

- What's inside a node when:
 - It has channels on multiple networks (different kinds of external names)?



What does that mean?



Inside vs. outside names

- Another way of distinguishing names
 - That all “belong” to the same node
- Names depend on your viewpoint...



“Outside” names

- Giving a name to the source or destination on a network layer
 - Source address to enable N:1
 - Destination address to enable 1:N
 - Same address to enable bidirectional communications

“Outside” names

- Names of a party
 - Node names
 - Interface names




Node vs. interface

- Node
 - Where processes run
- Interface
 - Network attachment point

Node vs. interface

- Node
 - Source/sink of all network channels at a single place
- Interface
 - End of one network channel



Node names


- Unique across
 - All nodes within a layer
- A node may have multiple
 - Node names
 - On the same or different layers
- Node names are equivalent
 - Within a node

Interface names

- Unique across
 - All endpoints within a layer
- A node may have multiple
 - Interfaces
- An interface may have multiple
 - Interface names
- Endpoint names are equivalent
 - Within an interface

Node name uniqueness

- Node vs. interface uniqueness




Strong vs. weak endpoint models

- We name interfaces AND nodes
 - What happens when we use those names?

Strong

- Names refer to the interface (channel end)
 - If a message arrives at a node from network A, it must be addressed to the endpoint address where that node attaches
 - All names belong (in effect) to the interface
 - Like the name of the doors of a house



Like at Downton Abbey

*Guests to the
front door*



**Tradesmen
around the back**




But they both end
up in the house



Weak

- Names refer to the node
 - Even if assigned to the interface
 - If a message arrives at a node, it can be addressed to any endpoint address where that node attaches
 - All names belong (in effect) to the node
 - Like the names of a house




As in, All Roads Lead To Rome



Kinds of outside names

- Ethernet
 - A name for a channel endpoint for Ethernet messages (Ethernet layer)
- IP
 - A name for a channel endpoint for IP messages (IP layer)
- TCP, UDP
 - A name within an IP endpoint called a port (we'll get back to that shortly...)

A look inside the endpoint...



“Inside” names

- Names within a party
 - A communication source or sink from the view within the endpoint



Inside names...

- What do we need to refer to?
 - The data itself (objects)
 - The process that uses or creates it

Object related names

- File names (static data)
 - C:\Users\guest\Desktop\file.doc
 - /usr/include/stdio.h
- I/O names (infinite source/sink of data)
 - LPT1:, COM0:
 - /dev/pty0, /dev/ttya, /dev/eth3
 - Socket descriptor (complex data structure)

Process related names

- Process
 - 8842
- Thread
 - 223
- Other related names
 - User – 521, “reihe”
 - Group – 9111, “lasr”

OS Review

- Process
 - Smallest independent running program with its own memory space
 - Resources include program code, memory, and thread(s)
- Thread
 - Smallest independently-schedulable running program

Why we prefer process names to...

- Thread names
 - Single address space of a process ensures each process name is unique
 - Thread names might be unique only within their parent process space
- File, I/O, etc. names
 - In this class, comm. endpoints are TMs
 - A TM more closely maps to a process

Properties of inside names

- Syntax
 - Defined only for that node
- Value
 - Unique within the node

Meaningless as network identifiers


Job of an OS

- Coordinate resource sharing
 - Share memory, CPU capacity, devices, channels, etc.
- Provide abstractions
 - Of machines
 - To allow multiprocessing
 - Of other resources
 - Like the network layers

How do OSes abstract layer endpoints?

- Socket
 - Created by ARPAnet research (RFC33, 1970)
 - A communication endpoint from the view of the “user” (program)
 - Usually two-way
 - Basically: a socket is an inside name for outside communication...

What's that mean?



We need to tell the computer's operating system to connect

Process A

To channel X

A socket is A's inside name for the outside name (channel X)

Room for confusion . . .

- Unix-style systems also use sockets for machine-internal IPC
 - Where one process communicates to another
 - With no actual (or even virtual) networking involved
- Our concern is with network sockets

Inside and outside


- How do we link:
inside names and outside names?

Linking the two

- Bind
 - Currently common OS convention
 - OS operation linking an internal I/O name to an external communication layer name

Two sides to a socket

- Server side
- Client side



A socket (either side)

- Ask the OS to create a placeholder
 - Attached to the process that creates it
 - A data structure that will link to the outside

```
if ((sockfd = socket(AF_INET, SOCK_STREAM, 0)) == -1) {  
    perror("Server: socket");  
}
```

- Now I've got a socket
 - But I need to attach it to an inside name

Common kinds of sockets

- Datagram (e.g., Ethernet, IP, UDP, ATM AAL0)
 - Direct to the channel
 - Separate messages
 - Individually addressed
- Stream (e.g., TCP, ATM AAL2-5)
 - Two-party association (“connection”)
 - Two steps:
 - Establish shared context with an address
 - Exchange data using that shared context
- Others are possible, but not common

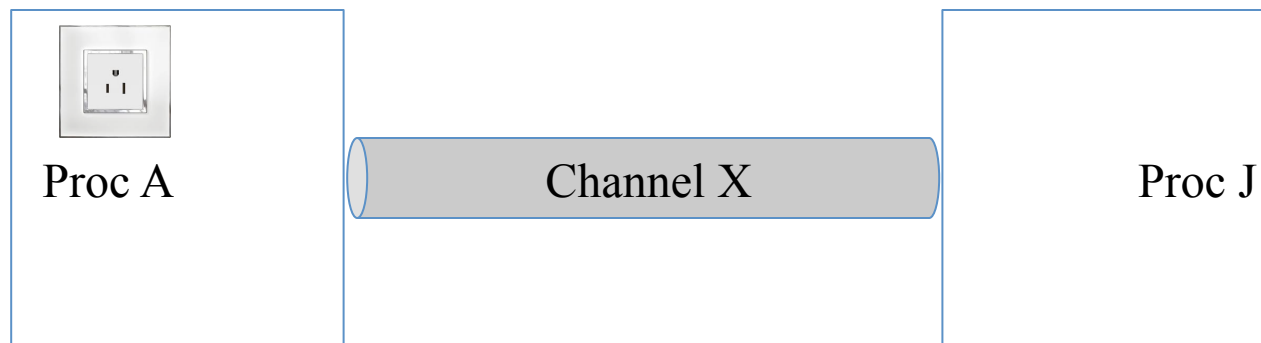
Bind

- Link a socket to an address on the server end
 - For TCP
 - Describes server end of the connection
 - For UDP
 - To limit messages you receive
 - To avoid source-addressing each message sent

```
if (bind(sockfd, (struct sockaddr *)&server, sockaddr_len) == -1) {  
    close(sockfd);  
    perror("Server: bind");  
}
```

Server first steps

- Socket
 - Create a channel placeholder local to the process
- Bind
 - Link the channel placeholder to an external name



Stateless: receiving messages

- Recvfrom
 - Accept a message
 - Indicate who it is from (other end)

```
recvlen = recvfrom(sockfd, inbuf, MAXDATASIZE, 0,  
                  (struct sockaddr *)&client, &clientlen);
```

Stateless: sending messages

- Sendto
 - Send a message
 - Indicate who it is to (other end)

```
if (sendto(sockfd, outbuf, outbuflen, 0,  
          (struct sockaddr *)&client, clientlen) < 0) {  
    perror("Server: sendto failed");  
}
```


Server side - connections

- Listen
- Accept

Listen

- Wait for incoming connection
 - Mark socket available for incoming requests
 - Prepare for someone to connect to the other end
 - Limit max waiting to be handled

```
if (listen(sockfd, MAX_CLIENTS) == -1) {  
    perror("Server: listen");  
    exit(1);  
}
```

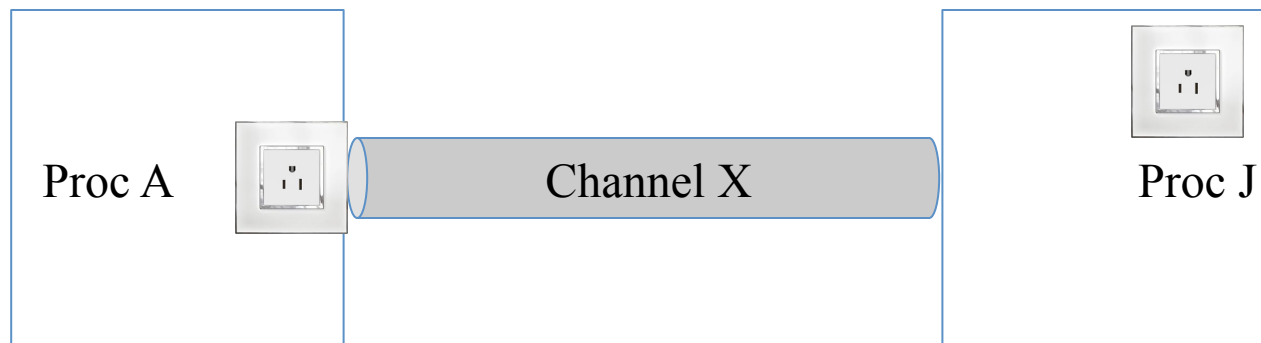
Accept

- Turns a socket into a socket pair
 - Socket pair defines a connection (both ends)
 - Now someone is connected to the other end
 - NB: in Unix, a socket and a socket pair are both described by the same data structure (a Unix socket)

```
new_fd = accept(sockfd,  
    (struct sockaddr *)&client,  
    (socklen_t *) &sockaddr_len);
```

Client side – connections

- Socket
 - Need something to connect to
- Connect
 - Connect socket to the channel



Connect

- Initiate a connection to a remote end
 - Indicate the remote end
 - Wait for the connection to be accepted

```
if ((connect(sockfd,  
    (struct sockaddr *)&server,sizeof(server))) == -1) {  
    perror("Client: connection error");  
    exit(-1);  
}
```

sockaddr and names

- What is the sockaddr?
- A data structure containing an external name
 - A name the client can use to specify which server socket to connect to
- In practice, an IP address and a port
 - Which is, remember, the type of name used by TCP and UDP

Client and server data exchange

- Send
- Recv

Send

- Write data on the connected socket
 - Same as `sendto` with `NULL` remote endpoint
 - Can be wrapped with a `write` call for simpler use

```
if (send(sendsock, sendbuf, strlen(sendbuf), 0) == -1){  
    perror("Client: send");  
    exit(1);  
}
```


Recv

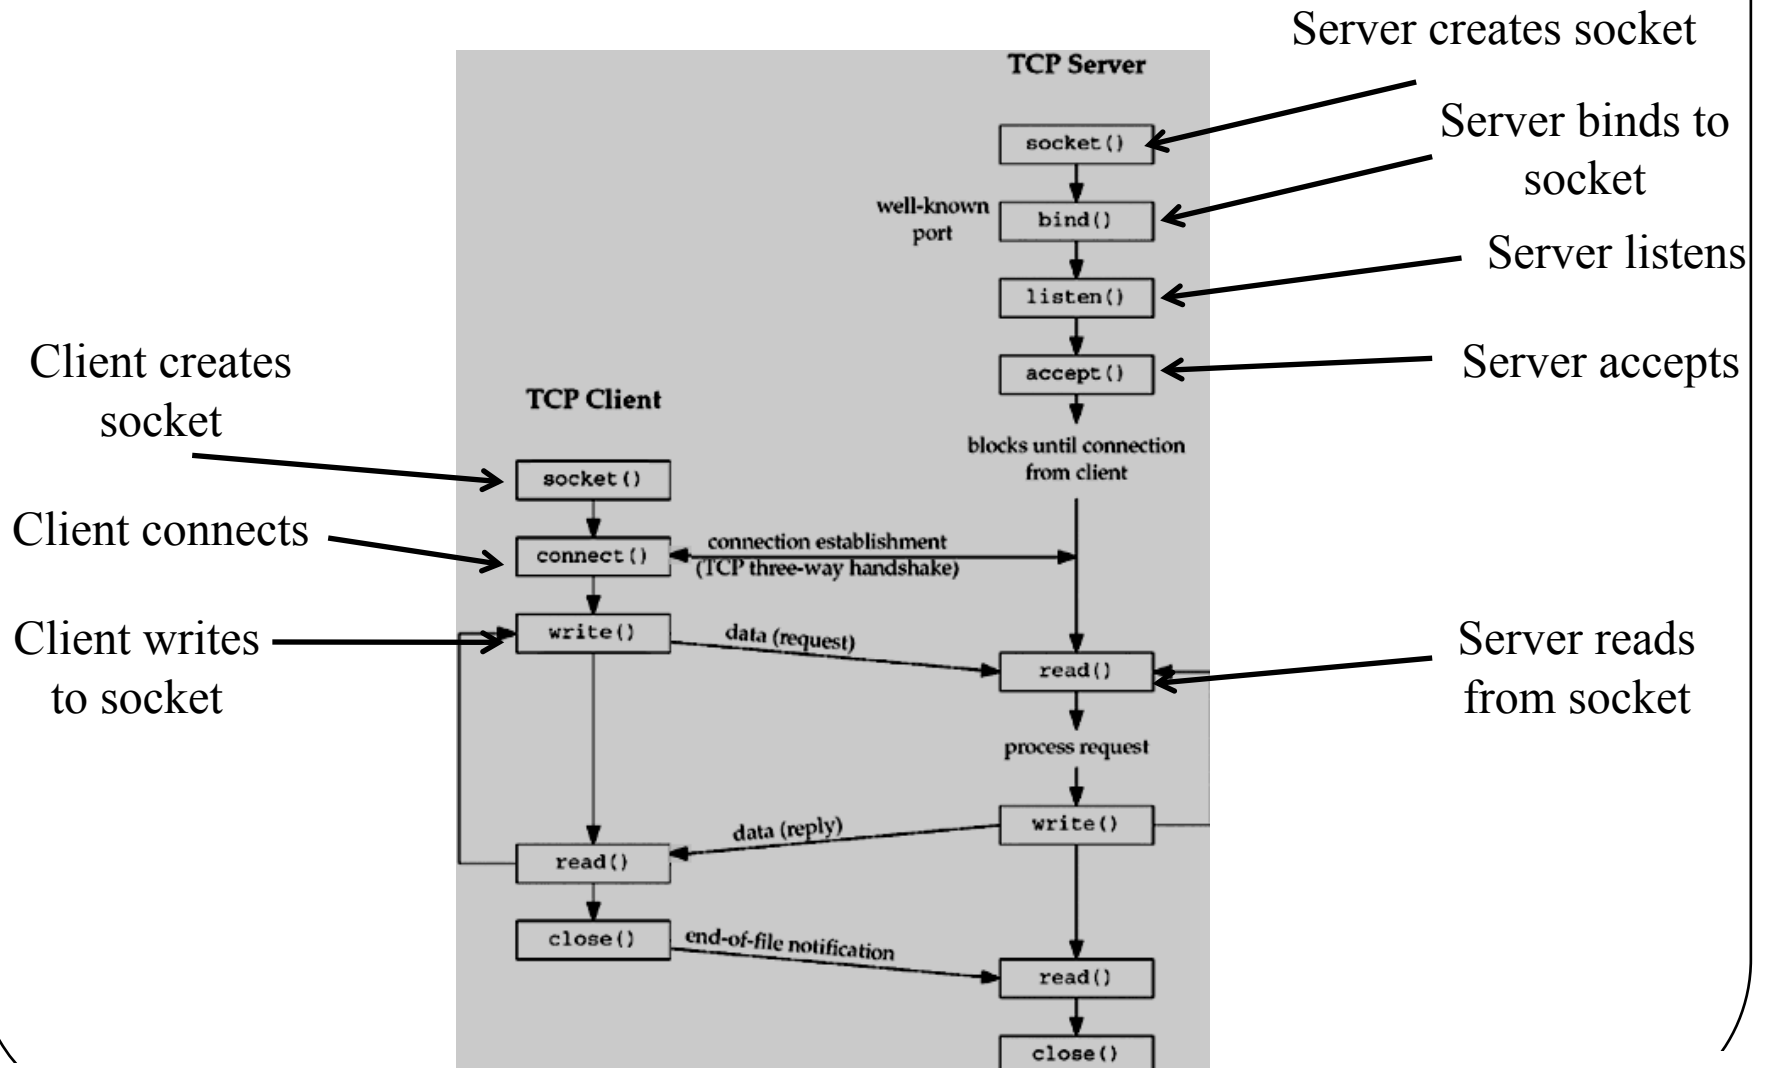
- Read data from a connected socket
 - Same as `recvfrom` with `NULL` remote endpoint
 - Can be wrapped with a `read` call for simpler use

```
if ((num = recv(recvsock, &buf, MAXLEN, 0)) == -1) {  
    perror("Server: recv failed");  
    exit(1);  
}
```

Putting it all together

- How do you arrange a client/server connection with sockets?
- Server creates a local `socket` and `binds` to it
- Client creates a local `socket` and `connects` it to the server's external name
- Server `listens` on the socket and `accepts` incoming messages
- Client `writes`, server `reads`

For example,



Issues

- Messages without bind
- A horse with no name
- Socket type and boundaries

No bind, no problem

- Stateless (connectionless) messages
 - Bind indicates local end
 - What if you omit bind?
 - The OS figures out where the message should go and adds the source address itself

Automatic source addressing

- What do you already indicate?
 - Destination address (required in sendto)
 - Includes “address family”
 - Unix-speak for layer name
 - Only one layer of each name!
 - Includes destination address
 - Use that with an internal (route) table to pick an outgoing interface
 - Set the source address to the outgoing interface

What about ports?

- Recall:
 - Port distinguishes different TCP or UDP layer endpoints within a IP layer
- How do you know the one to send to?
 - Someone tells you!
 - From a published list
 - Because you're replying to a message (or within a connection) already know

What's your port number?

- Messages
 - The one you know to send to
 - The one you got a message from (to reply)
- Connections
 - The one a server LISTENs on
 - The one a client CONNECTs to

Port numbers

- Ports are local to the pair communicating
 - Identifies the socket (thus the process) on each end
- Sometimes ports have common meaning
 - At “first contact”, they help you pick who you’re talking with (i.e., client-side)
 - That’s why they’re registered by IANA

Two meanings of ports

- During first contact – expected process
 - E.g., web server (80), secure server (443), email server (110), etc.
- After that, *just* an endpoint identifier
 - At the TCP/UDP layer

Port meaning

- By common convention (assumption)
 - Groups:
 - System ports (80, 110, 53)
 - User ports (8080, etc.)
 - Dynamic ports (unassigned!)
 - Assigned to “services” (TM expecting messages)
- By other coordination
 - Because you and the other endpoint agree
 - Port can mean anything you (and they) want

Having no name

- Bind with no name?
 - Technically, you cannot
 - Bind to “0” = ANY (i.e., “don’t care”)
 - Works for IP address, TCP/UDP port
- What happens when you need a name?
 - If you picked ANY, the OS assigns you one
 - Address = based on path, from ones you “own”
 - Port = pick one not in use

Socket types and boundaries

- Sometimes send/recv boundaries match
 - E.g., the channel preserves the boundaries
 - Sending messages
 - Sending data over a connection with markers
- Sometimes, not so much
 - E.g., TCP!
 - If you send 100 bytes, that might go in one TCP message, two, three, etc.
 - When the other side recvs, you don't know what data is ready

Summary

- Naming is more than just for networking
 - Names inside the machine
 - Binding between inside and outside names
- Names are linked in a set of steps
 - We used Unix as an example
- Names also set expectations
 - E.g., port number implies TM type (“service”)