

Automatic Naming

CS 118

Computer Network Fundamentals

Peter Reiher

Outline

- What is automatic naming?
- Why automatic?
- Designed-in
- Asking someone else
- Figuring it out for yourself
- Issues

What is automatic naming?

- Assigning a name to a network entity without human intervention
- Usually very dynamically
- Usually at the moment when it is first needed
- Often using different names for the same thing at different times

Why automatic?

- “Because it must be!”
- Ease of configuration
- Adapting to changes

Because it must be!

- Without a name, what can you do?
 - Anonymous reporting (N:1)
 - Broadcast announcements (1:N)

Not all that useful, but...
we can use these to get a name!

Ease of configuration

- Convenience matters
 - Plug-and-play, Zero-touch, etc.
- Complexity is painful
 - How many devices do you own?
 - Are they all configured the same way?
 - What if you had to configure them explicitly?

Adapting to changes

- Mobility
- Renaming

Mobility

- Change of physical location:
 - Changes network location
 - Topological or geographic names change
 - E.g., USC IP on campus, TimeWarner at home
 - Changes network
 - Name space changes
 - E.g., phone number on 4G,. IP address on WiFi

Renaming

- Change by the network operator
 - E.g., area code “split”
- Change by the user
 - E.g., off-campus WiFi then VPN to campus

How can you get a name?

What are the options?



Alternatives


- Design-in (preconfigure)
- Pick at random
- Ask someone else

Designed-in sub-options

- The \$1 solution
- Dude, where's my card?
- Getting the boot


The \$1 solution

- Maximum cost of globally unique names
 - Use a USD \$1 serial number as your name
 - Put the \$1 in the device (or whatever)



Ethernet

- **Two part solution:**
 - IEEE assigns OUI
 - Organizationally-Unique Identifier
 - \$2,575 per block of 16M addresses (2^{24})
 - \$0.0001535 per address (6,515 per \$1)
 - OUI assignee manages the block



Ethernet addresses

- All Ethernet devices have:
 - Fixed
 - Wired-in or write-only by manufacturer
 - Unique Burned-in (BIA) / hardware (EHA) address
 - Broadcast (all 1's)
 - Writeable
 - To change your BIA (to replace systems)
 - To add multicast addresses


POTS, non-SIM cellphones

- Assigned by a hierarchy of authorities
 - ITU country codes, country area codes, ...
 - POTS – paired to the “tail circuit” (house wire)
 - Non-SIM cell – paired to 7-byte MEID
(Mobile Equipment ID; 32-bit ESNs ran out in 2008)


44 + 28 + 9043-4310

country code of the country you are calling	area code of the area within the country you are calling—not all countries use an area code	local number
---	--	--------------

Dude, where's my card?



Full-size (FF)



Mini (2FF)



Micro (3FF)




Nano (4FF)

SIM-based cellphones

- GSM phones have two names
 - The phone (IMEI)
(International Mobile Equipment ID 14 digits, 6.228 bytes)
 - The SIM card (Subscriber Identity Module)
 - Includes a 20 digit ICCID (IC circuit ID)
- Telco links ICCID to your phone number
 - Also checks your IMEI isn't blacklisted (stolen)

Getting the boot

- Power-on configuration
 - Files on disk, USB, floppy
 - Flash memory
 - *PROM (EEPROM)
 - Ask the user (let's hope not . . .)



Figuring it out for yourself

- Pick me a winner!
- Parental support

Rolling the dice...

- If the number space is large enough
 - Why not just pick one?
 - What could go wrong?



People names

- Hierarchical in spirit
 - Given name(s) are “random”
 - But are they?
 - What if your last name is common?

Rank	Male name	Female name
1	Michael	Lisa
2	John	Mary
3	David	Susan
4	James	Karen
5	Robert	Linda
6	Mark	Donna
7	William	Patricia
8	Richard	Lori
9	Thomas	Sandra
10	Jeffrey	Cynthia
11	Kevin	Kimberly
12	Scott	Tammy
13	Joseph	Deborah
14	Steven	Pamela
15	Timothy	Brenda

IPv4 link local

- 169.254.x.x
 - EXCEPT first 256, last 256 (RFC 3927)
 - Based on MS Automatic Private IP Addressing (APIPA)
 - Pick randomly, do a test to confirm
 - Works only on the local link
 - Where the test works (ARP)
 - NEVER relayed
 - E.g., on your Ethernet

Pseudo-what?

- Random
 - Having no predictability
 - A sequence with maximum disorder
- Is a single number ever random?
 - No such thing!
 - Random applies to a sequence

Random number generation

- Cannot be generated by a TM in finite time
 - A TM would read only a finite tape
 - TM + finite tape = predictable output

So what do we do?

True random

- Need an external source of infinite entropy
 - A random physical event
 - E.g., radioactive decay, thermal noise, Brownian motion




Pseudorandom

- Deterministic, but appearing random
 - Unix rand()
 - Ethernet BIA
 - Disk access times
 - Keystroke delays
 - Mouse movements
 - Repeatable
 - Useful to replay simulations

“Spot” the difference



pseudorandom



raindrops

Eyeballs aren't always useful

2089986280348253421

1706798214808651328



2306647093844609550

5822317253594081284

8111745028410270193

8521105559644622948

Compute the difference



IPv6 link local

- FE80::/10
 - Assign based on MAC address
or
Pick randomly (RFC 4193)
 - Do a test to confirm
 - Works only on the local link
 - Where the test works (ND)
 - NEVER relayed

iOS Ethernet anonymity

- When configured
 - Every time device wakes from “sleep” (almost never, FWIW)
 - Pick a new random MAC
 - Hope it doesn’t collide (!)
 - There is no test!
 - Avoids “fingerprinting” SSID requests
 - Some stores monitor these

Asking DAD for help

- Duplicate Address Detection
 - Any general mechanism
 - “DAD” is specific to IPv6
- Works where?
 - IPv4: yes
 - IPv6: yes
 - Ethernet: NO

IPv4 duplicate detection

- Use ARP
 - Send an ARP probe for yourself
 - Source IP = none
 - Destination IP = broadcast
 - Owner MAC = yours (*presumed unique*)
 - Query for = the tested address
 - Do NOT send a query *from* the tested address
 - It will overwrite the cache of others!
 - Possibly even the existing owner!

Crossing the streams?



- ARP vs. IP
 - Different layers
 - IP nodes sit on both
 - Nodes on shared links
- Are these gateways?
 - Not quite
 - We never translate, only encapsulate (stack)

Implications for IPv4

- IPv4 addressing
 - Ask one network layer for help with another
 - Exchange ARP so IP can autonumber
 - Exchange ARP so IP can discover
 - IP on shared links doesn't exist alone!
- What about non-shared links?
 - Addresses are assigned statically

IPv6 DAD

- Use IPv6 Neighbor Solicitation
 - Same basic principle as IPv4
 - Ask to see if anyone has the desired address
 - If nobody asks, we get it

IPv6 Neighbor Solicitation

- IP-level replacement for ARP
 - But IPv6 has no broadcast
 - Use multicast instead
- How?
 - Could multicast to “all nodes” (like ARP does)
 - Instead multicast to MAC based on IPv6 addr
 - Only the node we want joins that group
 - NOBODY ELSE IS BOTHERED!

More parental support – IPv6

- Global IPv6 address
 - Listen for a Router Advertisement
(or ask routers via Router Solicitation)
- Create an address you know is unique
 - Combine RA information with Ethernet MAC
- Do a test to confirm
 - The test is only on the local link
 - Avoids MAC collisions
 - But the address is good globally
 - RA part is assumed unique

IPv6 example

- Listen for router advertisements
 - Collect them as they come in
- For each RA received on an interface
 - Combine the router prefix with the MAC BIA
 - Also join an IPv6 multicast based on the BIA

Asking someone else

- A horse with no name
- Name servers for self-namers

A horse with no name

- Asking a question without an ID
- Getting an answer without an ID?

Asking a question...

- How do you start?
 - If you don't know who to ask, broadcast the question
 - If you do know who to ask, send directly
- What's your address?
 - At the layer you need to know, NONE (typically “0”)

What layer do you ask?

- IPv4
 - Another layer (generally)
- IPv6
 - Your layer (always)

IPv4

- Mixing the layers
 - On a *different* layer that already has an address
 - E.g., broadcast Ethernet ARP with your MAC address
 - E.g., ATMARP request to LANE server on known circuit
- Same layer
 - IP (with UDP inside) to DHCP server
 - On the same layer to a server
 - Using source address 0

IPv6

- IP directly
 - Neighbor Discovery
 - Source address = 0

Getting an answer...

- Broadcast
 - When you didn't know who was asking
- Unicast
 - When you do
(e.g., when the request is over a different layer)

What can someone else tell you?

What are the options now?



What can someone else tell you?

- Just the facts
 - An address based on a table
- The facts and stuff
 - An address based on a table
 - A file that could have anything
- A loan
 - More specific information
 - Organized by type
 - Loaned out, then recovered for reuse

Reverse ARP

- ARP
 - Broadcasts request providing IP address
 - Owner replies with corresponding Ethernet MAC
- RARP
 - Broadcasts request providing Ethernet MAC
 - Server replies with corresponding IP address

RARP limitations

- Only provides an IP address
 - Systems often need more, e.g., default router, DNS server (to avoid bugging the roots), etc.
- Requires preconfigured server
 - Each expected request must match an entry
- Runs on its own protocol
 - Like ARP, this isn't over IP; it's over Ethernet

BOOTP


- Bootstrap Protocol
 - Still needs a static, preconfigured table
- Replacement for RARP
 - Runs over UDP over IP
(rather than Ethernet directly)
 - Also provides a file to retrieve
 - That file can be a script, a program, or a table

DHCP

- Dynamic Host Configuration Protocol
- Replacement for BOOTP
 - Runs over UDP over IP
 - Explicit way to manage specific configuration parameters
 - Managed via leases
 - Assignment has an expiration; can be renewed, released
 - Allows easy reassignment

Steps in DHCP

- ARP-like two-phase address assignment
 - Client broadcasts (IPv4) or multicasts (IPv6) a UDP DISCOVER request
 - DHCP servers *all* broadcast/multicast a UDP lease OFFER
 - Client picks one offer and *unicasts* a REQUEST
 - DHCP server *unicasts* a UDP ACK



Why two phases?


- Multiple servers can make an offer
 - Client picks only one
 - Second phase confirms selection
 - Offers are released after a time if not selected

B/Mcast vs unicast

- Unicast where possible
 - If you know which DHCP server you want
 - If you've already leased some info
 - E.g., and you go back to get more...


DHCP relay

- A little like proxy ARP
 - But in both directions



DHCP relay

- A little like proxy ARP
 - But in both directions



What can DHCP configure?

- DHCP offer
 - Information critical to configuring the channel
 - IP address
 - dynamic from a range or static based on a table
 - Default router
 - And “netmask” (indicates shared link addresses)
 - Lease time
 - DNS server
 - To avoid root overload
- DHCP inform
 - Other additional context
 - Time server
 - Network Time Protocol
 - Web proxy
 - Address, parameters, etc. for shared caching
 - Just about anything else

DHCP events

- Request
 - Client searching for initial offers
- Offer
 - Servers making initial offers
- Request
 - Client picking one offer
- ACK
 - Server confirming offer
- Renew
 - Client asking for lease extension
- Release
 - Client asking for lease cancellation

USB

- Master (host), assigns to slaves
 - Assigned each time a device is plugged-in
 - 127 addresses (7 bits, 0=not set yet)

The single master controls “the world”

Name service for self-namers

- Recall: bind
 - Maps a process to a TCP/UDP port
 - How does another party find that port?
 - It knows the number (IANA list, pre-agreement)
 - It knows the name, but not the number
- Register your name
 - Contact the DNS that has your name:IP map
 - Add the portname:portnum entry too

Issues

- Telling everyone else
- Configuring DHCP
- Impact to communication in progress

Telling everyone else...

- How do others know your new name?
 - Esp. if you make one up
- Remember the DNS?
 - Can also map persistent names to changing ones
 - lever.cs.ucla.edu -> IPv4 address that isn't 131.179.192.136
 - IMAP@lever.cs.ucla.edu -> port that isn't 110

Using the net to find names

- Remember the need for glue?
 - DHCP’s “glue” to the client:
 - Router address
 - Even better when it’s a “default” router
 - Channel subnet mask
 - What’s reachable without contacting the router
 - DNS server
 - A way to get names without needing a default router
 - It needs to be reachable either on the shared channel or via the router indicated

Configuring DHCP server

- DHCP makes leases
 - Where does it get its land (resources)?
- Currently:
 - Manual configuration
- Experimentally:
 - Another server (“Dynamic DHCP Configuration”)

Pros and cons

- Design-in (preconfigure)
 - Pro: easiest, known to work
 - Con: won't deal with mobility, changes
- Pick at random
 - Pro: second easiest, might work
 - Con: might not (verify?), finding others is hard
- Ask someone else
 - Pro: easy for the client, allows coordination
 - Con: right back where you started for the server!

Impact on in-progress comm.

- What happens to connections or relays using addresses that change?
 - Continue using the old name
 - How do you know if this is even possible?
 - Shift to the new name
 - What if there isn't one?
 - What if there's more than one?

Summary

- Giving a name to yourself can be easy
 - Verification is needed
 - Using that name beyond the shared link is harder
- Most naming involves
 - Assumed uniqueness
 - Asking someone else
- Getting started is still manual
 - True “zero configuration” is very rare