

Layer Optimization

CS 118



Computer Network Fundamentals

Peter Reiher

Where are we at?

- We understand communications over direct channels
- We understand building networking from layering and relaying
- We understand how a DAG explains layering
- We understand how to build routing for relaying purposes movements through a DAG using recursion

For example,



HTTP->TCP

TCP->IP

IP->802.11

802.11->physical

Now we need to
relay through C

A sends an HTTP
request to E

We recurse down
the DAG

```
graph LR; A((A)) --- B((B)); A --- C((C)); C --- D((D)); D --- E((E));
```

And back

--


[illegible]

IP->ATM

ATM->physical

Now relay through D

Relaying through D




IP←-ATM

IP->ethernet

ATM<-physical

ethernet->physical




D now uses
its DAG

And back
down

Now relay
to E

Delivering at E




HTTP←-TCP

TCP←-IP

IP←-ethernet

ethernet←-physical



E now uses
its DAG

And the message is
delivered to E's web
server

Where are we and where next?

- So now we know how to use networking between multiple network points
- Being able to communicate at all is more important than anything else
- But other things are important, too
- Like performance, reliability, security, and other properties
- How can we optimize the basic networking to achieve these goals?

Outline

- Background
- Deficiencies
- Performance
- Emulation
- Examples...

Networks and optimizations

- Using fully general mechanisms can be expensive
- Common special cases can be optimized to reduce costs
- Optimizations are possible at many semantic levels of networking

But where to optimize?

- Is optimization a layer or a communication issue?
- Where do we optimize?
 - At some layer?
 - Or for some connection?

Intra-layer vs. intra-communication?

- Optimizations involve shared context
 - Layers sharing common mechanisms
 - Connections managing shared state
- Either one can support optimization
 - Connections coordinate explicitly
 - Layer members coordinate implicitly

If layers, which ones?

- Do optimizations occur only at certain layers?
 - No
- Are optimizations typical at certain layers?
 - Yes, for several reasons
- Do optimizations interact across layers?
 - Absolutely

Living at a layer

- Optimizations can occur at any layer
 - They're increasingly used at many layers

Living at a layer

- Some occur more often at certain layers
 - Most information errors are at the physical layer
 - Once corrected at the next layer up, they tend not to occur again
 - So optimizations based on these errors often at low levels

Living at a layer

- Some optimizations occur because of a layer
 - I.e., TCP provides an ordered data stream but IP does not
 - So TCP corrects ordering, but IP does not
 - Therefore, any ordering optimization occurs above IP

What connection?

- Optimizations often share state over a connection
- State can be
 - hard (maintained)
- or
 - soft (recoverable)
- State can be for one connection or a group
- State can be explicit or inferred

What connection?

- Not everything is associated with an explicit, stateful connection
 - I.e., there's more than TCP, web, and e-mail
caching?

Now let's explore how to optimize:

- Deficiencies

end to end principle

- put application specific network functionalities at the endpoints, not in the middle
- functionality in the middle should be generally usable by all applications
- not a hard and fast rule
- a principle that is influential

- Performance

- Emulation

Deficiencies

- Optimizations sometimes overcome deficiencies in the communication
- For example, deficiencies in:
 - Integrity behave in a way that has a property that it does not have
 - Authentication be sure that the receiver does this
 - Privacy receive the content of the message

Why do you care about deficiencies?

- Deficiency impedes communication
 - You can't share state, which means you can't share information
- Deficiency impedes relaying
 - If you can't relay for others, they can't communicate
- Deficiency impedes networking
 - The two above also mean you can't automatically manage your network configuration, routing, DAGs, etc.

Integrity

- Definition
- Types:
 - Corruption
 - Loss
 - Tampering

Integrity

the quality of being whole

- Definition: the quality of being whole
 - WYGIWWS: what you got is what was sent
- For a message:
 - Not split up
 - Not missing pieces
 - Not altered (accidentally or deliberately)



Integrity: Corruption

- Accidental alteration
- For a message:
 - Symbol changed (noise)
 - Symbol is ambiguous (equivocation)
 - A portion is deleted



Integrity: Loss

related to corruption

- Missing everything!
 - Degenerate case of corruption
 - Receiver doesn't know a message arrived
- Why?
 - Not sent (by origin or relay)
 - Not received (by destination or relay)
 - Corrupt beyond recognition



It's 2:20pm, do you know where you are?

- Time is fleeting...
 - But it keeps coming up
 - How do you detect loss?
- Useful to know about time
 - Longest time until delivery




Integrity: Corruption vs. loss

- Corruption
 - A message arrives
 - Unrecoverable error
- Loss
 - No message arrives
 - Timer implies loss
- Indistinguishable when:
 - Destination name is corrupted or missing
 - Source name is corrupted or missing
 - Key portions of message are corrupt or missing
- Difference
 - How much of the message is “gone”

Integrity: Tampering

- Deliberate alteration
 - To corrupt
 - To alter to different content (thought to be non-corrupt)
- How?
 - Intercept and retransmit (e.g., during relay)
 - Overlap physical signals



Integrity: Corruption vs. tampering

- Corruption
 - During origination, receipt, or relay
 - Detect via error checks
- Tampering
 - During relay or receipt
 - Detect via integrity checks
 - That a relay can't 'fake'
- Difference
 - Intent
 - Probability of generating different but valid message
 - Tampering is similar to the worst case for corruption

What about order?

- Order is not a property of a message
 - It is a typical property of a channel
 - Only a deficiency if you need to assume it
 - We'll come to that when we talk about channel emulation
- misordered delivery when i thought it was all delivered

How much integrity protection?

- Easy to encapsulate
 - You don't want to modify contents anyway
- Relay might become more difficult
 - May need to change portions of the message
 - E.g., hopcount, route path record, etc.
 - Does the integrity cover those?
- So you might not protect the entire message

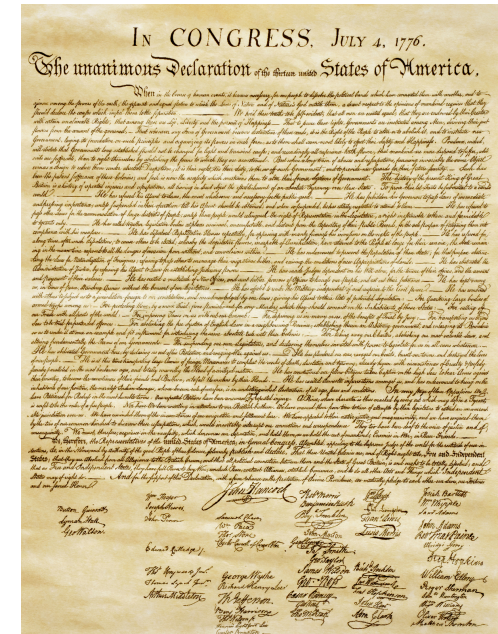
Authentication

- Definition
- Types
 - Origin and/or destination
 - Control
 - Content

Authentication

you can identify who created the information

- Definition: ensuring that particular information was created by a particular party
- For a message, ensuring:
 - All important elements of the message
 - Were created by the sender



Authentication: Name

- The source and/or destination noted in the message comes from the source that generated it
 - The message says it came from John, and it actually came from John
 - The message says it goes to Ben and that name came from John too
- Why?
 - Protects the entire message's path
 - Protects the endpoint machines

A stylized, handwritten signature in black ink that reads "John Hancock". The signature is written in a cursive style with a large, looping initial "J" and a circular emblem at the end.

Authentication: Control

- The message includes control signals that come from the source that generated it
 - Such as what layers it uses
 - And parameters to those layers
 - Ensuring that sender used those layers and those parameters
- Why?
 - Layers use state machines
 - Protects operation of the state machines



Authentication: Content

- The content of the message comes from the source that generated it
 - I.e., the data that is shared with John actually came from John

protect integrity without providing authenticity

if you want to know who sent it, then you need authenticity



- Why?
 - Protects information that the machines share

Why bother with all three?

- Why do we separate:

- Identity
- Control
- Content

three of these may have different parties sign the followig
we go up and down layers at the relays
we may want to authenticate what happens at the relays

- They could be signed by different parties

- Different endpoints, different layers, etc.



Why?

- For example, content comes from the “top” of the stack
 - So the top layer should authenticate content
- The identity might be associated with a proxy at a lower layer
 - So that layer should authenticate identity
- The control might be for a very low layer
 - So control must be authenticated there

How much to authenticate?

the message will change from place to place

- Same problem as integrity
 - Easy to attest to the source of the entire message when encapsulated
 - Hard to make that guarantee if portions change
 - Again, might want to authenticate only part

Privacy

- Definition
- Types
 - The origin and/or destination
 - Control
 - Content

Privacy

- Definition: hiding the information in a message from all parties except the receiver
- For a message:
 - Hide the origin
 - Hide the message
 - From everyone but the receiver



Privacy: Identifier

- Hide the endpoints

- Who sent it
- Who receives it

metadata: who communicate to whom

clever data mining: find out a lot about who talked to who, when, and perhaps for how long?

- From whom?

- Source: everyone except receiver
- Destination: everyone except relay and receiver
 - Perhaps even from relay

- Why?

- “who talks to whom” exposes information!

Privacy: Control


- Hide the state machine control signals
- From whom?
 - Everyone except receiver
- Why?
 - Exposes what the state machine is doing
 - That information can be used to attack the machine
 - Or deduce things about the communications

Privacy: Content

- Hide the information shared between the parties
- From whom?
 - Everyone except receiver
- Why?
 - (should be obvious)

Impact of deficiencies

- Need to balance
 - Relay, source, receiver perspective
 - Various preferences, requirements, and limits
 - Various costs (time, space, CPU effort)



Performance deficiencies

- Time
- Space
- Energy

Time

a deficiency



- Rate
 - Messages per time

how much info from A to B in a given point of time
- Latency/jitter
 - Time per message (between send and receive)

it takes some time from getting one bit here to there

jitter - derivative of the latency; is it constant?
better to have constant latency rather than changing latency

Rate

- How many messages can you send?
 - Messages per unit time
 - $1/(\text{time to send a msg}) * (\# \text{ msgs sent concurrently})$
- How to improve?
 - Less time for each message (higher BW)
 - More messages sent concurrently (parallelism)
 - Messages at the rate the receiver supports
 - Messages at the rate the network allows

Flow control

- Messages arrive at the receiver's rate
 - Avoid overwhelming the receiver
 - Avoid using excess storage resources
- How?
 - Control pacing (inter-message timing)
 - Control number of unanswered messages
 - Use feedback from the receiver

Congestion control

- Messages arrive at the relay's rate
 - Avoid overwhelming the network
 - An aggregate, network variant of flow control
- How?
 - Similar mechanism as flow control
 - Different source of feedback (net, not receiver)

Latency

- How long for a message to arrive?
 - Time per message between send and receive
- How to improve?
 - Decrease distance between sender/receiver
 - Increase BW
 - ...

Space

- How much space to represent a message?
 - Bits per message
- How to improve?
 - Compress (remove predictable patterns)
 - Within a message, across messages, etc.
 - Encode efficiently

Energy

- CPU capacity
- Actual energy

CPU

- How much work to process a message?
 - How many opcodes?
- How to improve?
 - Save reusable results / avoid duplicate effort
 - Alternate algorithm
 - Same result a different, “cheaper” way
 - Different, “cheaper” result with similar properties

Actual energy

- How much energy to process a message?
 - Electrical power, heat to dissipate, etc.
 - Not just “green”; saves \$, heat, and space
- How to improve?
 - Reuse rather than recompute
 - Lower clock rates
 - Avoid conversions

Emulation

- Wires
- Boxes and bundles
- Transactions and beyond

Wires

- Making a circuit from packets
- Pseudowire



Circuits from packets

- Reliable info stream from messages
 - Ordered and reliable
 - Typically relies on endpoint state
 - Not necessarily guaranteeing performance
 - More *like* a wire than a message; not equivalent
- Examples
 - TCP from IP (Internet)
 - AAL 1-4 from ATM (ATM)
 - TP4 from CLNP (OSI)

Pseudowires

- A channel from messages
 - Ordered, reliable, static capacity and delay
 - I.e., performance emulation, too
 - As close to a channel as possible
- Examples
 - SONET
 - TDMoIP (TDM emulation over IP)
 - PWE (pseudowire emulation)

Order

- Circuits and pseudowires emulate channels
 - Most channels assume ordered signal transfer
 - Need to detect and correct misordering
- Examples
 - TCP over IP (Internet)
 - TP4 from CLNP
 - Not ATM! (ATM is never misordered!)

Boxes and bundles

- Boundaries
- Flows



Boundaries

- Marking edges between items
 - Multiple items in one message
 - An item that spans multiple messages
- Examples
 - IP message vs. its fragments
 - DCCP, SCTP
 - HTTP over TCP (stream)

Flows

- Grouping separate connections to act together
 - Striping (increased capacity)
 - Coordinated management (shared control)
 - Alternate/backup (fault tolerance)
- Examples
 - ISDN channel bonding
 - TCP control block sharing
 - Multipath TCP, SCTP

Transactions and beyond

- Transactions
- Translation
- Other services

Transactions and beyond

- Extend service beyond information sharing
 - Support specific structured interactions
- Why?
 - Mostly software reuse
(any of these can be built on any communication service)

Transactions

- Conditional information flow
 - Serial: send B only if A is complete
 - Conjunction (AND): send C only after A and B
 - Disjunction (OR): send C only after A or B
- Many variations:
 - N of M: send Z only if at least 3 of 6 alternates
 - Send Z only if exactly 3 of 6 alternates

Translation

- Convert one message to another
 - Occurs within the recursive block
 - Also occurs for “gateway” relays
- Examples
 - Language translation (content)
 - Format conversion (HTML to ASCII)
 - Display conversion (desktop to mobile web)

Other services

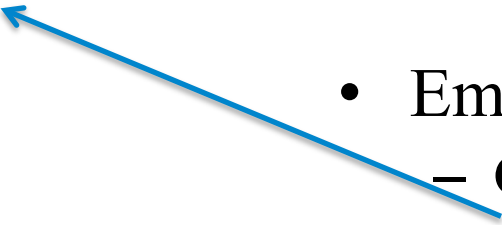
- If you can dream it, you can do it!
 - Any capability another user/system wants



- Some dreams are nightmares, though...



Overview of issues - sources

- Deficiencies
 - Integrity
 - Error
 - Loss
 - **Reordering?**
 - **Tampering**
 - **Authentication**
 - **Privacy**
 - Performance
 - Time
 - Rate (flow, congestion)
 - Latency
 - Space
 - Compression
 - Caching
 - Energy
 - CPU, actual energy
 - Emulation
 - Circuit/wire
 - **Reordering?**
 - Boundaries
 - Flows
 - Transactions
 - Translation
- 

Summary

- Deficiencies need to be fixed first
 - Can't communicate if we can't communicate
- Then performance should be addressed
 - Go fast, go cheap
- Then emulation
 - Make it look like the user wants it to look