

# Layer Optimization: Handling Loss

## CS 118

### Computer Network Fundamentals

#### Peter Reiher

# Intralayer examples

- Integrity
- Time

# Integrity

- Error
  - Detect accidental alteration
- Loss
  - Detect missing packet(s)
- Reordering (related mechanism)

# Detect or correct?

- If we can correct, there's no problem
- Here we focus on integrity failure
  - Based on detecting uncorrectable errors

can only detect, not correct

# What do errors look like?

- Errors in destination address
  - The message ends up elsewhere
- Errors in source address
  - The message affects the wrong FSM
- Errors in contents
  - Impacts other layers (up to the user)  
other layers, the AT NLP

It's important to know when you see an error

# Making errors visible

- Consistency checks
  - Portion of a message has a valid range
  - Value is outside that range (invalid values)  
sets a valid range for consistency within that range
- Redundancy checks
  - Add redundancy in a message
  - If error only alters one of the redundant values, the other indicates an error  
redundancy to see if error alters one of the redundant values

# Consistency checks

- Some values of some fields for some layers aren't allowed
- Examples:
  - Padding field in TCP header must be all zeros
  - Can't have unroutable IP address in Internet
  - IPv4 IHL field can't be less than 5

# Redundancy checks

- **Mathematical constraints**
  - Treat contents as numbers
  - Use only values that satisfy an equation
    - Matching:  $A = B$
    - Parity:  $(\sum a_i) \bmod 2 = p$
    - Checksums
    - Cyclic Redundancy Check (CRC)



# Example – IPv4 checksum

- 16-bit ones complement sum
  - Consider IPv4 header as set of 16 bit numbers
  - Add the numbers and hold the carry
  - Add the carry back in (carry-wraparound)
- Easy to implement in SW or HW
- IPv6 doesn't have header checksum  
only IPv4 has header checksum

# TCP checksum

- Independent of IPv4 checksum
- Taken over entire packet (except a few IP header fields)
- Basic method like IPv4 checksum
- Somewhat different computation method if running over IPv6
  - Interlayer dependency in this optimization . . .

# Loss

- Detecting when a message should have arrived but did not
  - Message completely lost
    - Never sent
    - Never arrives
  - Message source/destination error
    - Message arrives but cannot reach the intended receiver (or relay) FSM

# One way to lose a message

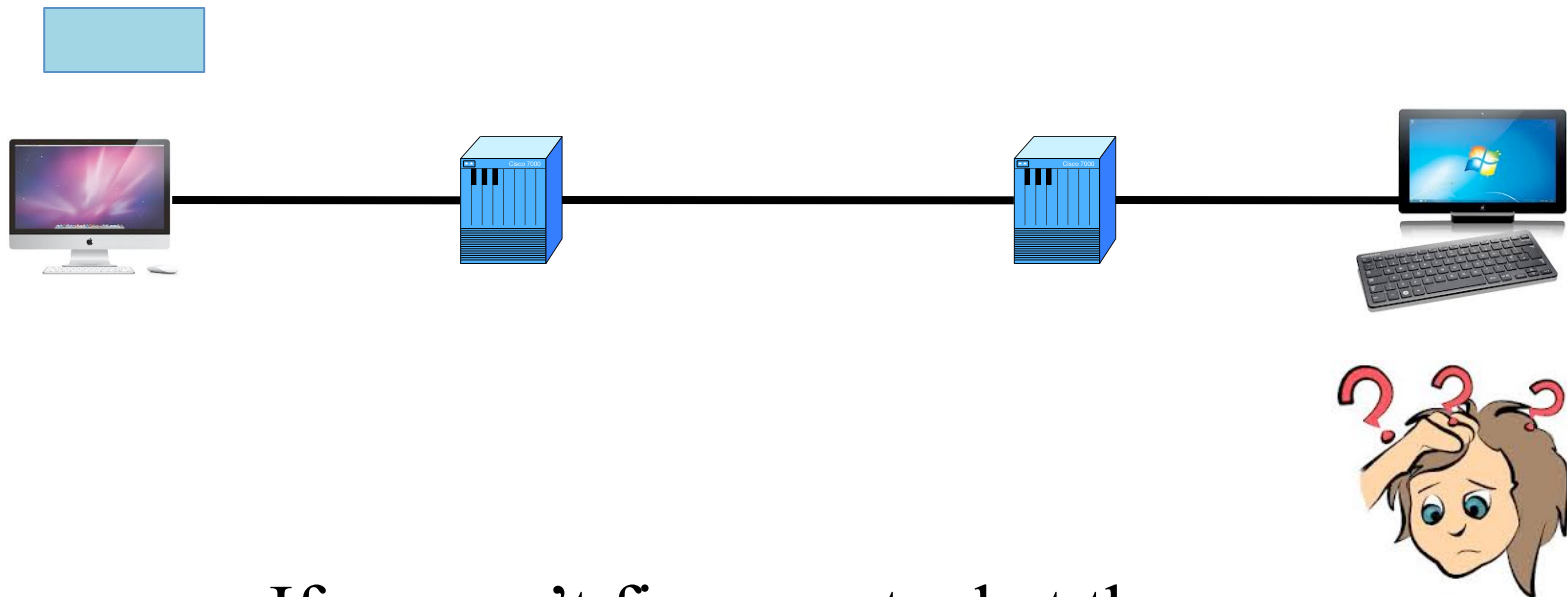
## *Timeout*



Different Ways to Lose a Message  
timeout  
error detection / correction  
flow control  
congestion control

# Another way to lose a message

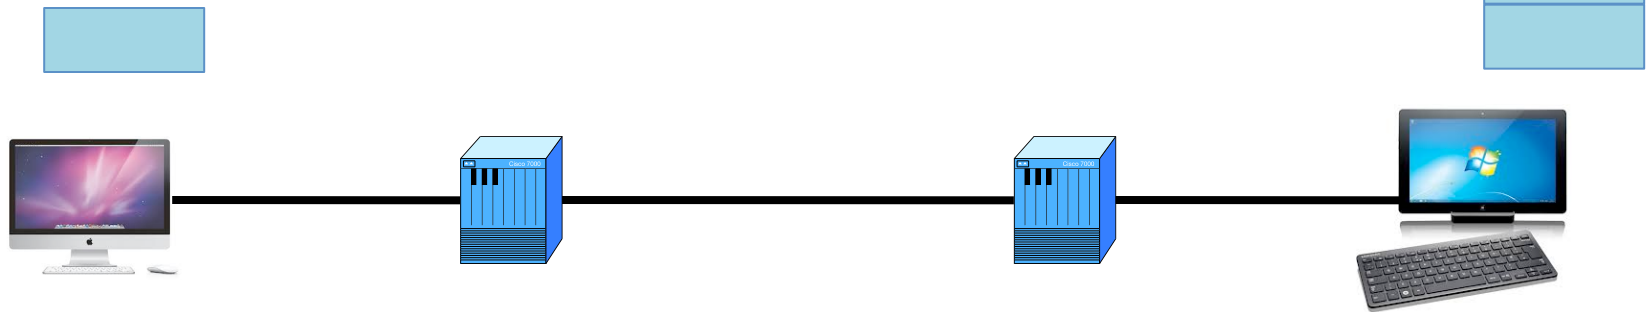
## *Error Detection/Correction*



If we can't figure out what the message is, it's as good as lost

# A third way to lose a message

## *Flow control*

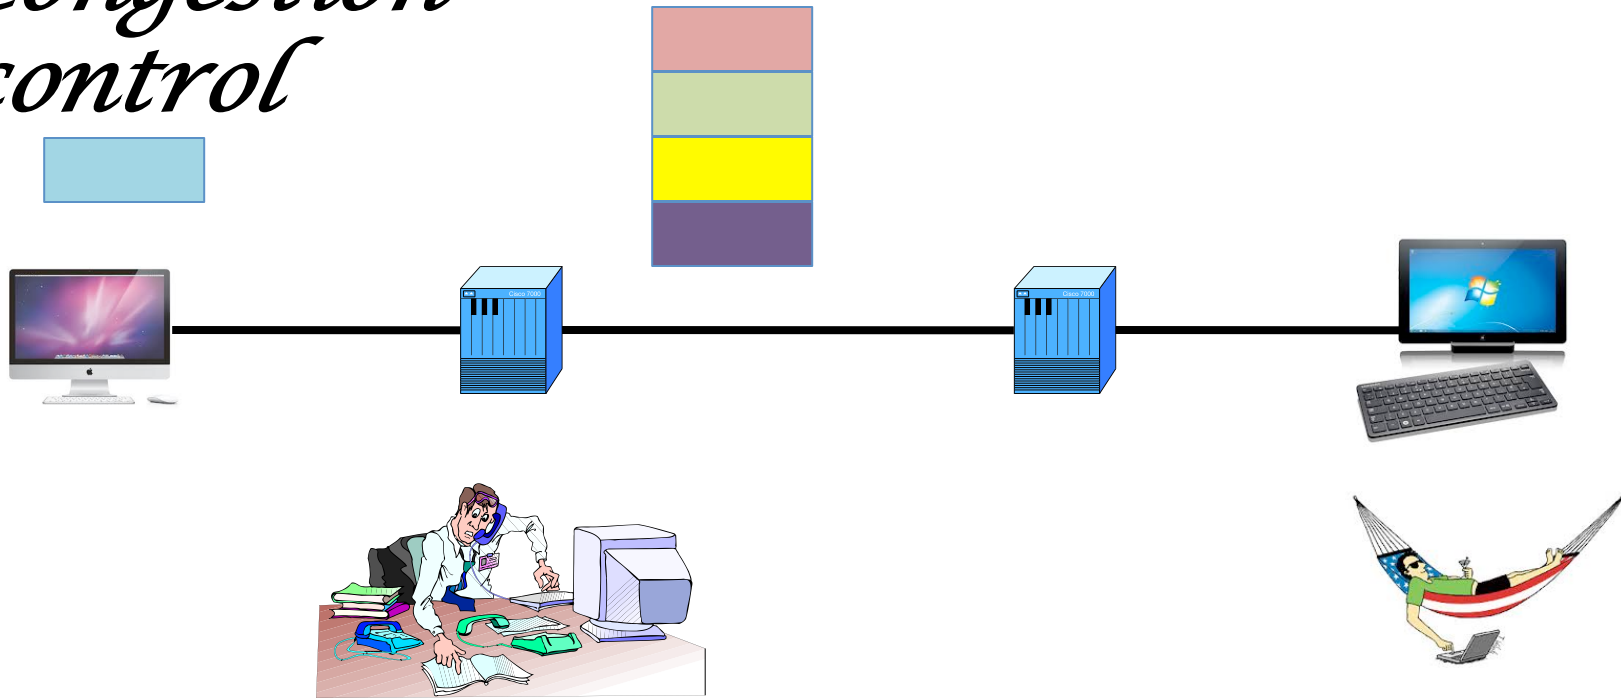


If the receiver is too busy and  
has nowhere to put it, he'll  
have to drop it



# A fourth way to lose a message

## *Congestion control*



If the network is too congested, it may be dropped

Even if the receiver isn't overloaded

# Detecting loss

- We've talked about this before
  - Need to keep track of time
  - So let's say we do...





# The key issue: what time?

- Different possible times:
  - Time expected
  - Latest time we want to wait
  - Time since last message sent
  - Time since last message received
- Let's say we know (or pick) one
  - Still need to know the value to use

# Estimating time

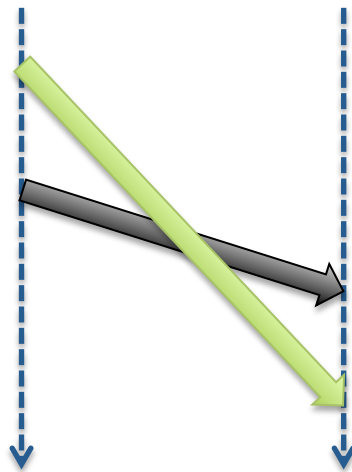
- Based on a network property
  - Number of hops (diameter)
  - Longest transmission time
  - Longest time held per relay \* # relays
- Based on past measurements
  - How do we measure time?
  - How do we aggregate measurements?

# Measuring time

- Need a timer
  - Can't wrap-around too quickly
  - Good resolution (time of smallest increment)
- Need to timestamp messages
  - Measure differences

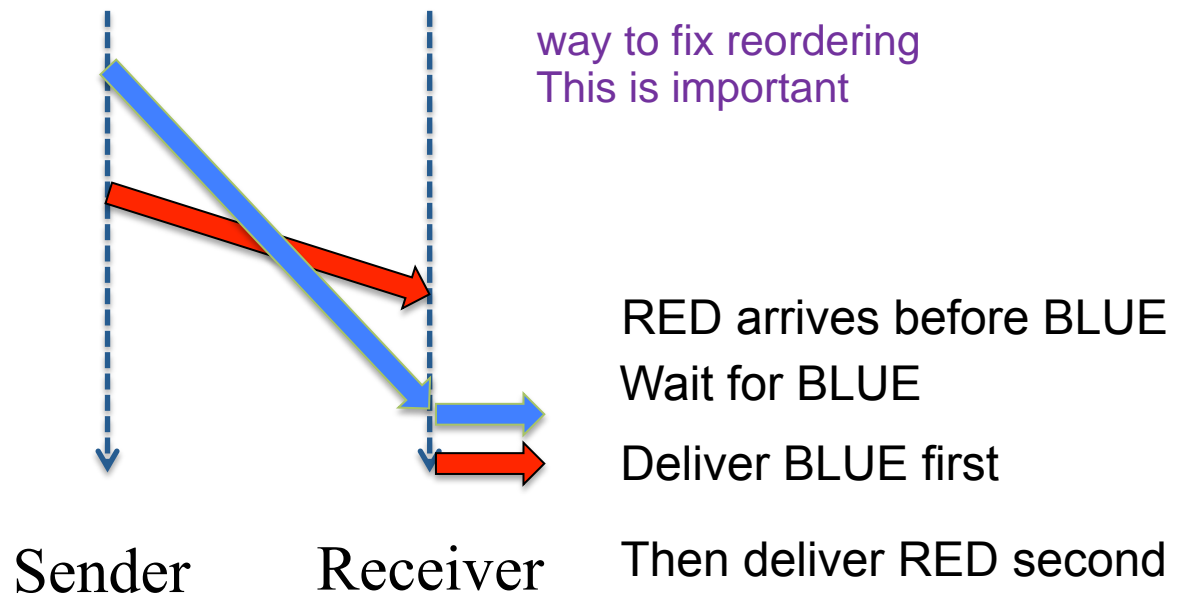
# Reordering

- Messages arrive without loss
  - But out of order



# Fixing reordering

- Hold onto messages
  - Keep the ones that come too early
  - Process once sequence gaps are filled



# What's hard about reordering?

- Need more state – a reordering buffer takes up space
- How much space?
  - Enough to store maximum displacement
- Where?
  - At the receiver, where things arrive too early
- Do you just wait for late messages?
- Or request a resend?
  - If so, how aggressively?

# TCP receive window

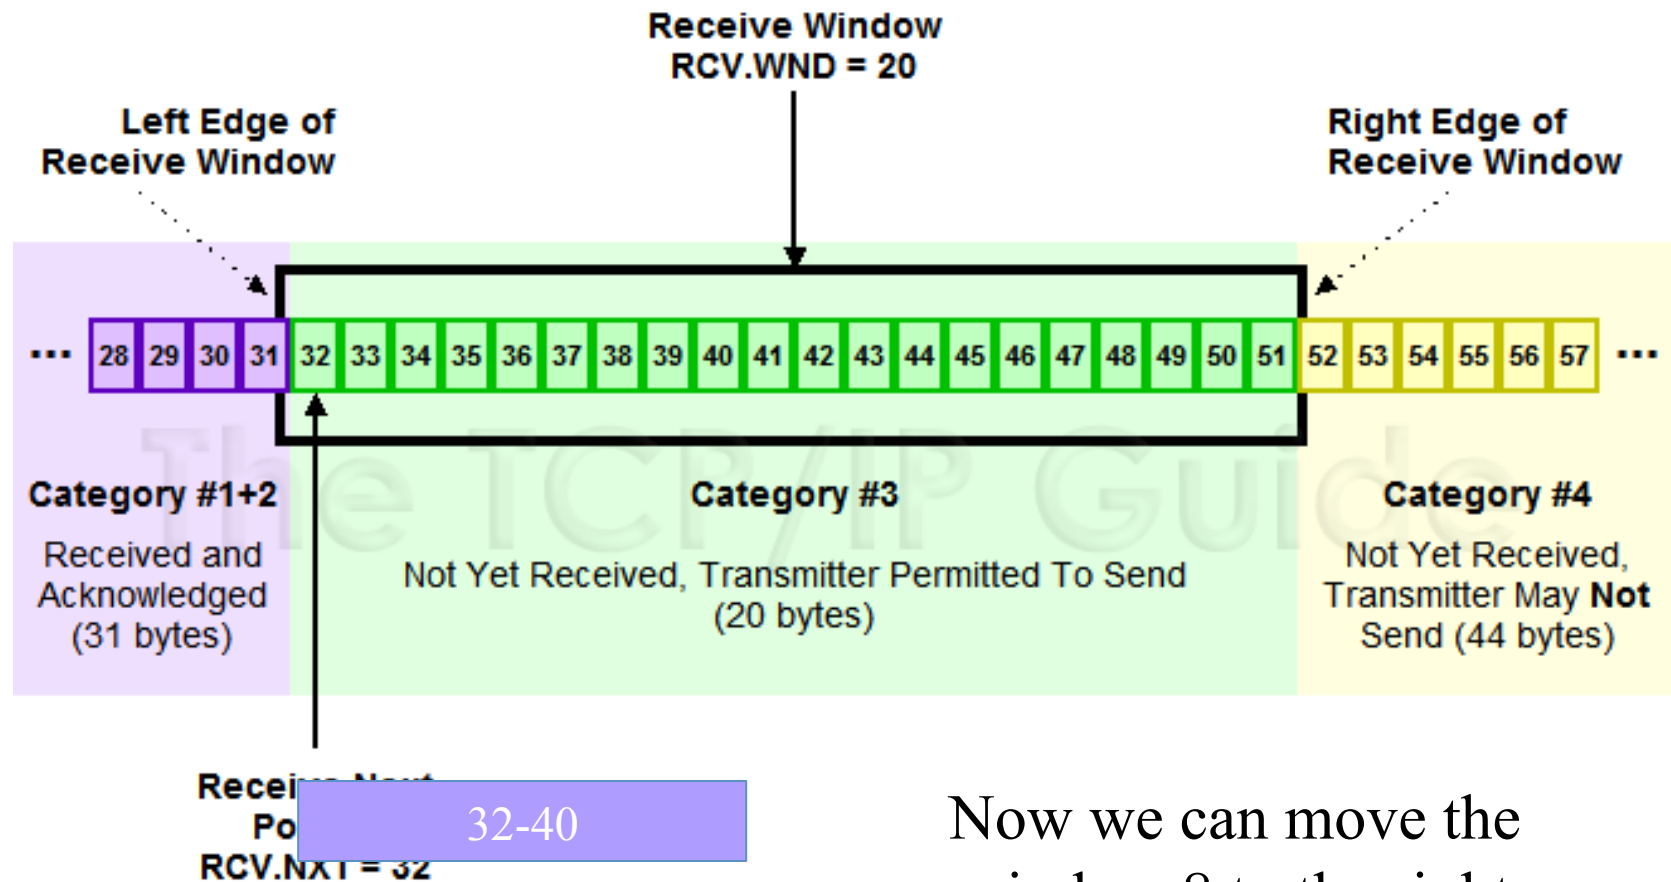
- TCP segments have a sequence number
  - Number indicates byte offset of each segment
- TCP receive window enables reordering
  - Make it large to handle large displacements
  - Left side = smallest offset not yet here
  - Right side = largest offset that can be held

# Receive window management

- Message arrives
  - Put it in the buffer if it fits
  - Buffer size =  $\text{RCV.WND}$
- $\text{RCV.NXT}$  (left side)
  - Move right to deliver info to the user (upper)
  - Stop and wait at the first gap
- $\text{RCV.NXT} + \text{RCV.WND}$  (right side)
  - Moves right (allow higher offsets) whenever left side moves right

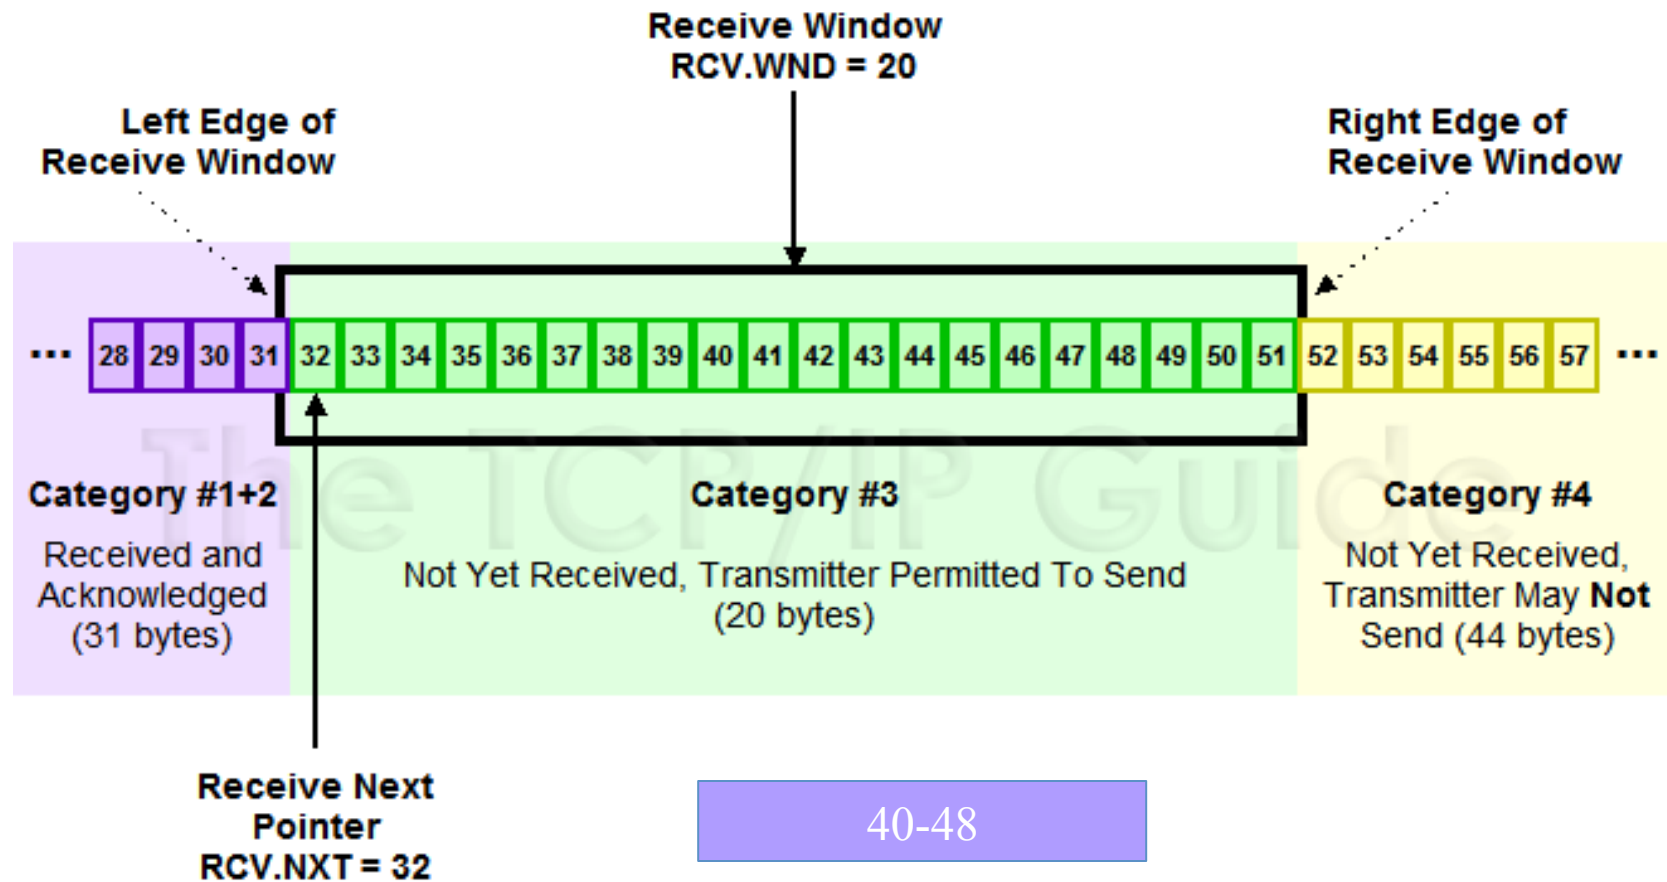


# Receive window



Now we can move the window 8 to the right

# Receive window with misordering



Can't move the window yet

# Hmm – what's that about the sender?

- That figure shows how the receive buffer is tied to the send buffer
- The receive buffer is for reordering
  - So what's the send buffer for?
  - Why (and how) are the two related?

# Time

- Flow control
- Congestion Control
- Latency management

# Flow control I

- Not overrunning the receiver
    - The receiver always can handle one message
    - Send that,  
wait for confirmation,  
send the next, ...
  - Why does the sender need to wait?
    - Message was lost
    - Message is waiting at the receiver
      - If receiver can handle messages as fast as the sender, not an issue
    - Solution: send one at a time
      - One message outstanding at any given time
- avoid waiting, send one message at a time

# Stop and Go

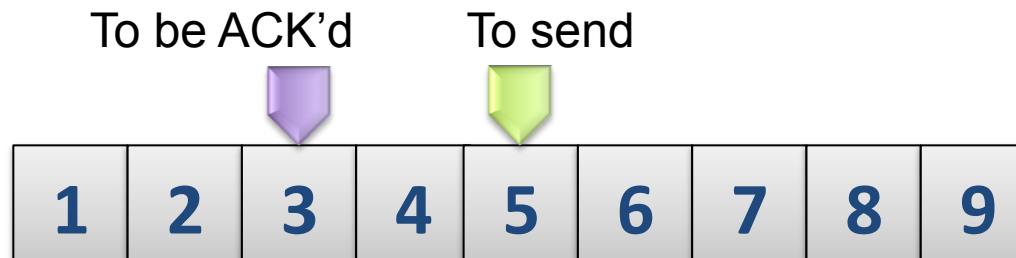
- Assume a sequence number
  - One number per message
- Sender and receiver work in lockstep
  - Both “walk” the number space
  - Both have “inchworm” behavior

# Let's take a walk

- Messages are numbered

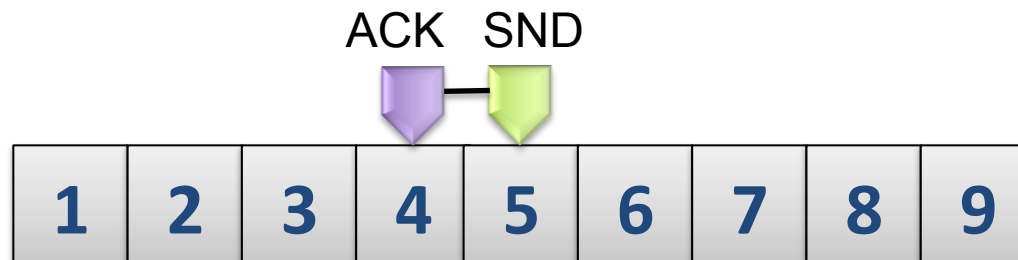


- Sender “walks” the line via receiver ACKs



# Limiting the walk


- **Send and receiver ACK linked by a limit**
  - $SND - ACK \leq N$  equation for linking send and receive
  - For stop-and-go,  $N=1$





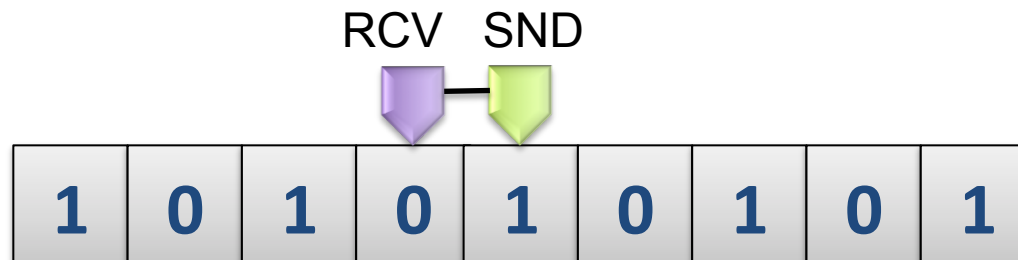
# A look at the exchanges

- One message per round trip
  - ACK indicates received and ready for next



# A look at the numbering

- If SND and ACK differ by at most 1, we don't need to number 1..999
  - OK to just number 0,1,0,1,0,1
  - Also known as “alternating bit” protocol



# Why do anything else?

- Receiver might be faster than sender
  - In which case it could handle more messages
- Learning that receiver has handled a message takes time
  - At least time to get acknowledgement to sender
- During that time, channel is not in use, receiver is not busy, sender is not busy
  - Lots of wasted resources

# What if we have more than 1 outstanding message?

- Messages could arrive out of order
  - As we've already discussed
- A message could arrive while receiver is busy handling an earlier message
  - Not possible with 1 message window Stop and Go
- If we allow multiple outstanding messages, we must handle both problems

# Flow control II

- How to handle these problems?
- As discussed, use a buffer to handle misordered messages
  - The receive window indicates max disorder
  - Also max “outstanding” held messages
- Hmm – let’s use that buffer two ways!
  - If receiver is busy when message arrives, put it in the buffer
  - Even if it is the next message to be received
- So we can send more than one at a time...

# Go Back N

- Assume a sequence number
  - One number per message
- Sender and receiver work in lockstep
  - Both “walk” the number space
  - Both have “inchworm” behavior


Just like stop-and go, but with N messages

# Let's take another walk

- Messages are numbered



- Sender “walks” the line via receiver ACKs



# Limiting the walk

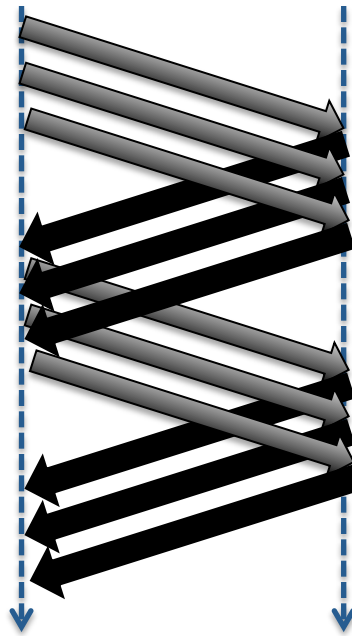
- Send and receiver ACK linked by a limit
  - $SND - ACK \leq N$
  - For Go-back-N,  $N > 1$





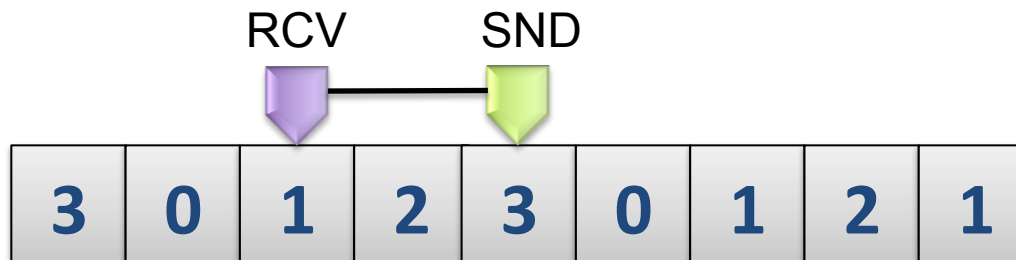
# A look at the exchanges

- N messages per round trip
  - ACKs indicates received and ready for next



# A look at the numbering

- If SND and ACK differ by at most  $N$ , we don't need to number 1..999
  - OK to just number  $0, 1, 2, 3, \dots, 2N-1$



# Why $2N$ values?

- $N$  outstanding values
  - Each RTT, the window can slide forward by  $N$
  - Need to prevent overlap from one RTT to next

# How big is N?

- How many messages before getting ACK?
  - Once you get the ACK for the first, you can send  $N+1$
  - ACK provides a “clock” to the pipeline
    - Every ACK/ $N+1$  pair acts like stop-and-go
    - Go-back- $N$  is like  $N$  overlapping stop-and-go

$N$  overlapping stop and go is a good analogy for stop and go

# About the receive window

- What if the receiver isn't fast enough?
  - Info (message) has to go into the buffer as fast as it arrives (or we have other problems!)
  - If the FSM doesn't release the info to the upper layer as fast as it comes in, there's a delay

# Recall receive rules

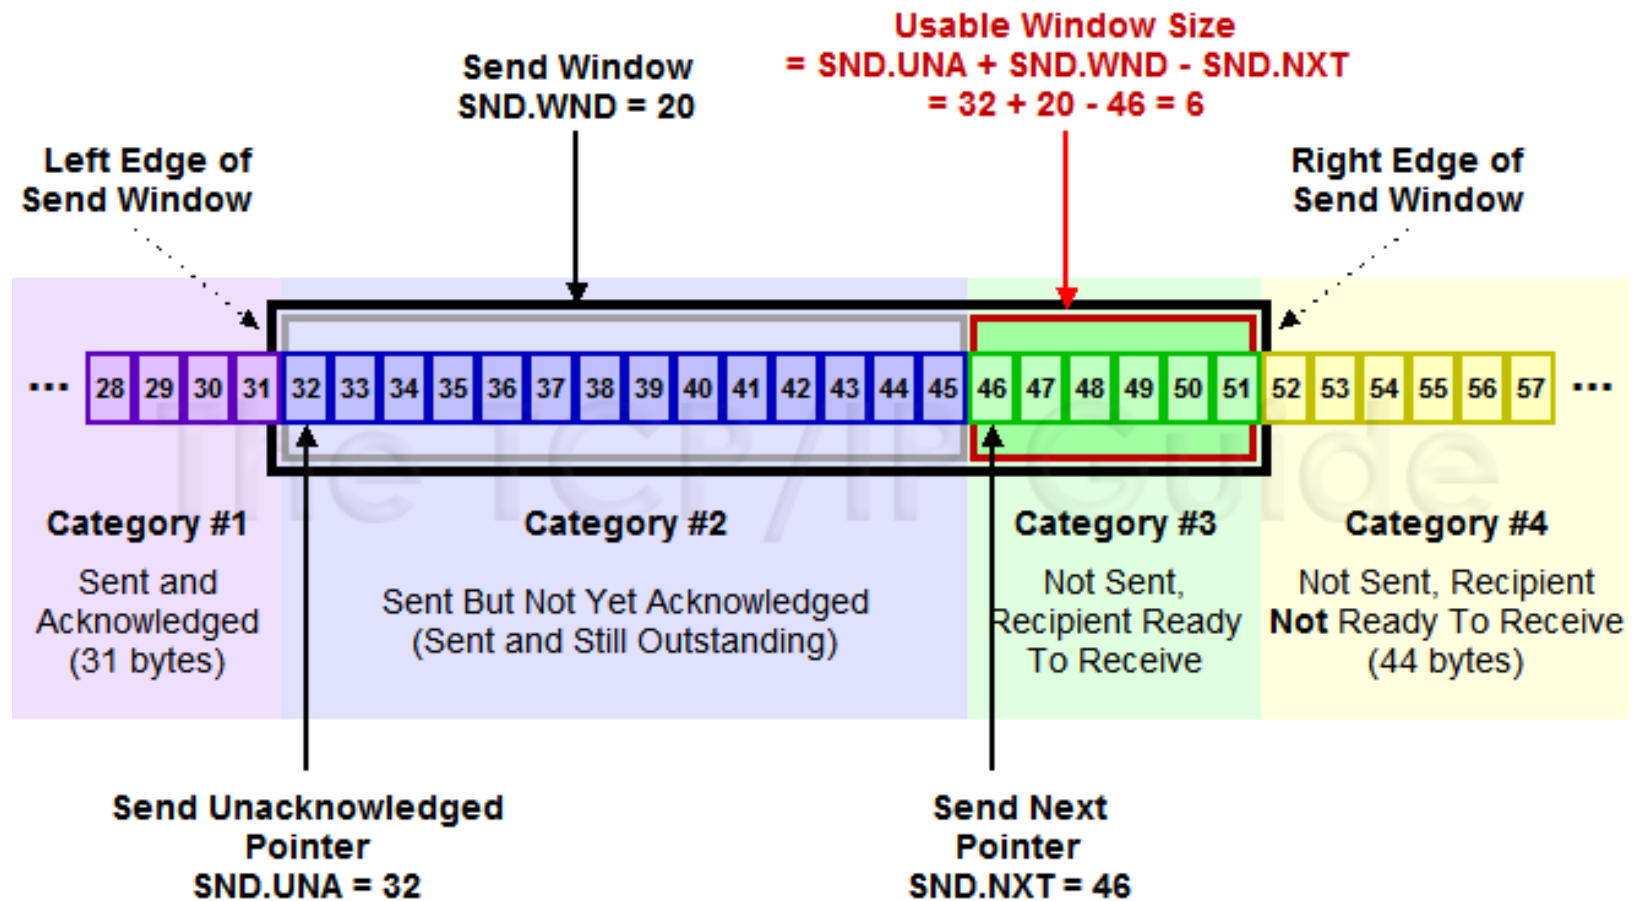
- Info (message) arrives
  - Place message in sequence
  - Move left side to the right until a gap
    - Pass that info to the next layer up
  - Right side moves to the right at the same time
- If the FSM is fast enough
  - The left side doesn't move immediately
    - Takes time – time to process the message

# Left and right

- If left side doesn't move, right doesn't
  - I.e., receiver isn't ready for new offset info
- How do we coordinate with the sender?
  - Sender has a similar buffer
  - $SND.WIN = RCV.WIN$ 
    - If smaller, we could have sent messages that could have been held by receiver – wasted resources
    - If bigger, we won't be able to send it anyway (we can only fill the receive buffer!)

- flow control - uses a window --> at sender, permitted to send out certain amount of data to the receiver; receiver checks what they have seen
- deal with misordered messages
  - receiver is capable of holding a certain amount of data until his FSM is capable of catching up with the different layers; ensure no overflow

# Coordinating SND and RCV



guarantee that we can send you all the data we have



# Combining loss + windowing

- Positive feedback (ACK)
  - Indicate what was received
- Negative feedback (NACK)
  - Indicate what is missing but expected
  - Always a gap after the last msg received

use NACK when we know something is missing

Use this info to coordinate retransmission

if negative ack, we can go ahead and do a retransmission

# Loss / windowing variants

- Stop and Go    one at a time request for retransmit
  - On timeout, send retransmit request
  - Only one message to ever request
- Go back N    back up to lowest ACK every trip
  - On timeout, ACK lowest missing sequence number
  - Sender “backs up” to where ACK indicates
  - Every round trip, backs up to the lowest gap
- Selective ACK    ACK everything you send ; selective repeat
  - ACK everything you get, ask to fill in the holes
  - Sender fills in only the holes

resend one packet at a time

# Congestion control

- Receiver might be ready, but is the net?
  - Don't want to overwhelm the network

control the congestion of data

when sender and receiver are both okay, but network is too busy

- We have some windows
  - Send = how much info can be outstanding
  - Recv = how much info can be reordered
- *Can* isn't the same as *should*

How much SHOULD be outstanding?

# Solution: congestion window

receive window stays fixed

- Receive window
  - Stays fixed (no benefit to adjusting)
  - As large as reordering max
  - As large as send pipelining too
- Send window
  - No larger than the reordering max
  - As large as is needed to keep up with the receiver
  - Not so large that messages are lost in the net
- OK, how big is that? 

increase the send window to keep up with the receiver