

# Layer Optimization: Security and Privacy

CS 118

Computer Network Fundamentals

Peter Reiher


# Another type of layer deficiency

- Some layers of protocol do not provide any security or privacy
- What if we want to have better security and privacy?
- What do we do to get them?

# What do we mean by “security?”

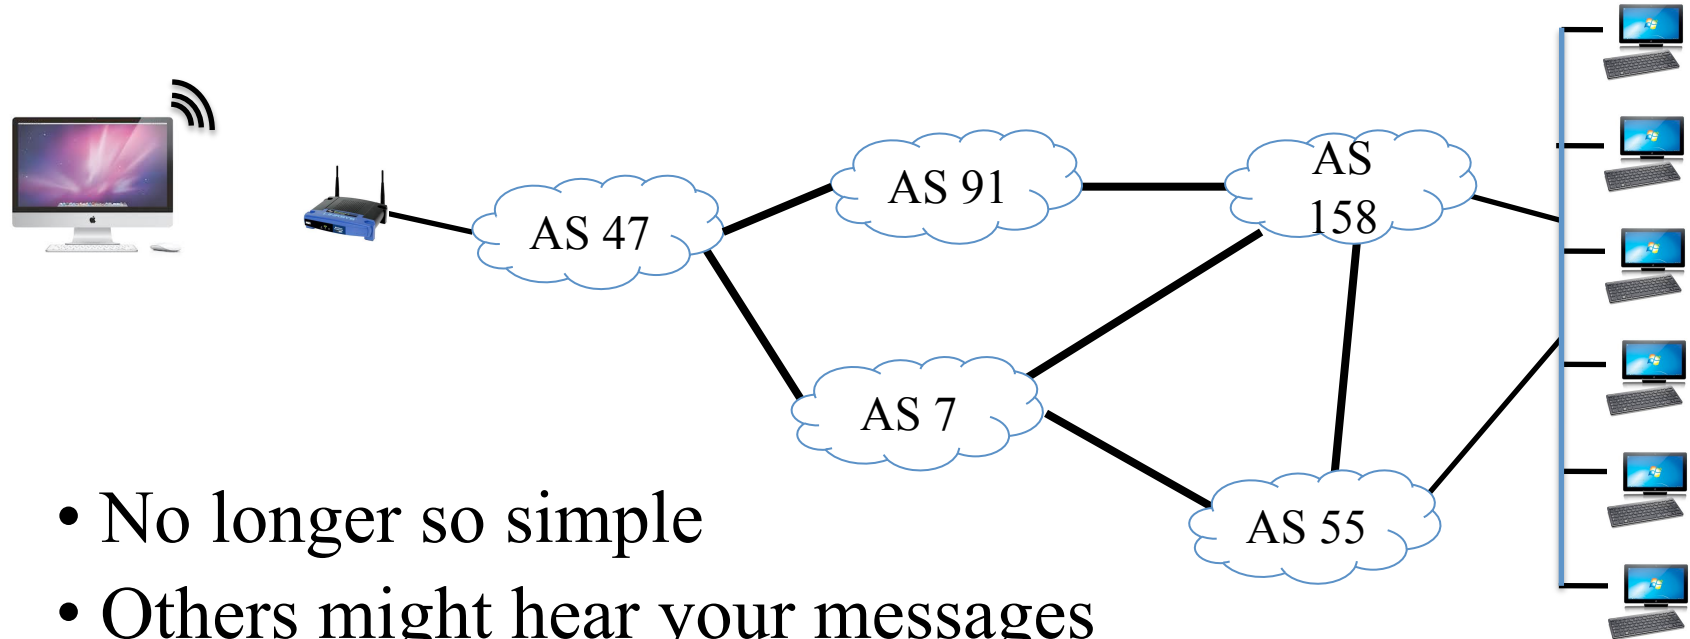
- Informally, providing some of three properties:
  - Confidentiality
  - Integrity
  - Availability
- In the face of adversaries attempting to compromise those properties

# Security on a single link



- A relatively simple problem
  - If this picture is accurate
- Nobody else can hear your messages
  - So confidentiality is good
- Nobody else can alter your messages
  - So integrity is good
- Nobody else can interfere with your messages
  - So availability is good

# Security in a complex network



- No longer so simple
- Others might hear your messages
  - So confidentiality is bad
- Others might alter your messages
  - So integrity is bad
- Others might interfere with your messages
  - So availability is bad

# Authentication

- Proving that something was created by a particular party
- E.g., a message was created by the user who appears to have sent it
- Vital property to achieve many security goals
- Since sometimes you will do things for some parties, but not others
  - Only works out if you can tell who is who

# Security

- Background
- Information protection
- Resource protection

# Background

- Basic mechanisms
- Key management



# Basic mechanisms

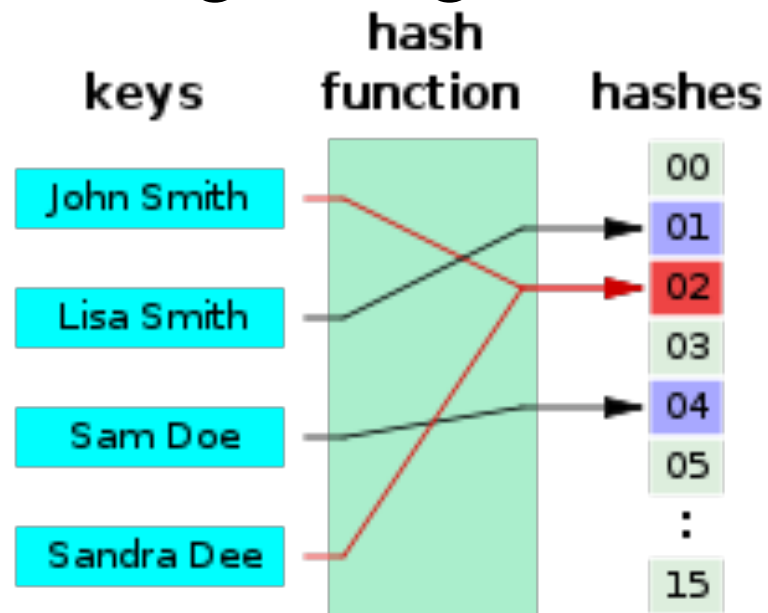
- For networks, primarily based on data manipulations
  - Hashes
  - Ciphers and codes
  - Signatures
- Also need to protect network resources
  - Need different mechanisms for that

# Security by data manipulation

- Put (or alter) data in packets to improve security
- Hashing
  - Integrity (detect tampering)
- Encryption
  - Confidentiality (obscure semantics/meaning)
- Signature
  - Authentication (identify source)

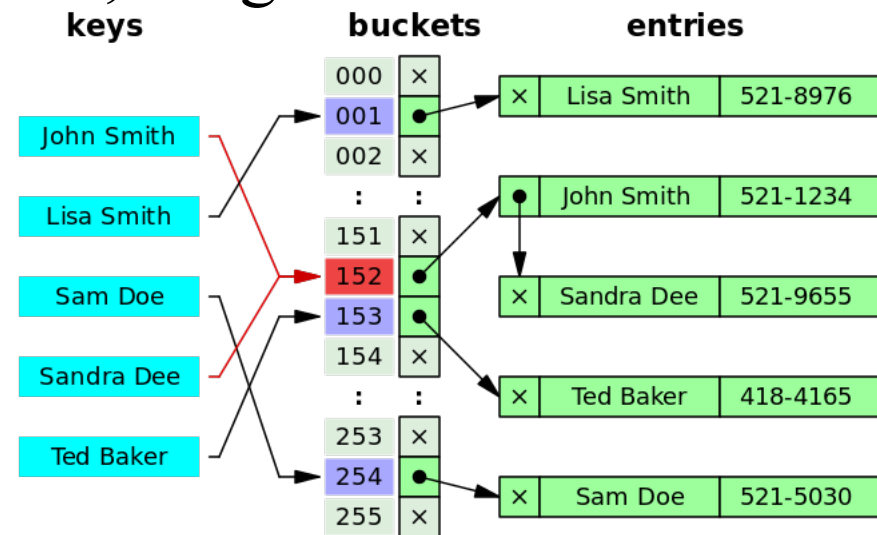
# Hash functions

- Maps a variable-length message onto a fixed-length “digest”



# Why hash?

- Scramble
  - Nearby messages yield very different digests
  - Distributes / scatters load, usage
  - E.g., hash tables
    - Rapid lookup
    - Avoids bunch-up
    - Collisions OK



# Cryptographic hash

- A hash function useful for cryptography
  - Scrambles and spreads (like any hash)
  - Difficult to “game”
- Anti-“game” properties
  - Unidirectional (non-invertible)
    - Difficult to generate another input for a given output
  - Rare collisions
    - Difficult to generate two inputs with the same output

# Why not just use a checksum?

- Checksums don't protect against tampering
  - Easy to generate a new message with the same checksum
  - Easy to generate two messages with the same checksum
- Cryptographic hashes do protect
  - Unidirectional and rare collisions make the above difficult

# Example hash functions

- Message Digest 5 (MD5)
  - The 5<sup>th</sup> attempt at a message digest
    - MD, MD2, MD3, MD4, MD5, now MD6
  - Weak – found a birthday attack
- Secure Hash Algorithm (SHA)
  - SHA-1 is weak
  - SHA-2 and SHA-3 well regarded
  - US Government designed

# What do you do with a hash?

- Publish it as-is
  - A “fingerprint” to validate as “untampered”
  - Assumes the published hash wasn’t tampered
- Use it in other algorithms
  - HMAC
  - Digital signatures



# Fingerprint checks

- E.g., for GPG software

d065be185f5bac8ea07b210ab7756e79b83b63d4  
091e69ec1ce3f0032e6b135e4da561e8d46d20a7  
fb541b8685b78541c9b2fadb026787f535863b4a  
5503f7faa0a0e84450838706a67621546241ca50  
d0cf40cc42ce057d7d747908ec21a973a423a508  
dc03ae4e4c3e8fe0583b37dd6c3124f94246d2f8  
4997951ab058788de48b989013668eb3df1e6939  
9456e7b64db9df8360a1407a38c8c958da80bbf1  
86fe0436f3c8c394d32e142ee410a9f9560173fb  
7cf0545955ce414044bb99b871d324753dd7b2e5  
01e62c45435496ff0e011255fb0ac1879a3bc177  
8dd7711a4de117994fe2d45879ef8a9900d50f6a  
9eb07bcceeb986c7b6dbce8a18b82a2c344b50ce  
a7a7d1432db9edad2783ea1bce761a8106464165

gnupg-2.0.27.tar.bz2  
gnupg-2.1.3.tar.bz2  
gnupg-w32-2.1.1\_20141216.exe  
gnupg-1.4.19.tar.bz2  
gnupg-1.4.19.tar.gz  
gnupg-w32cli-1.4.19.exe  
libgpg-error-1.19.tar.bz2  
libgcrypt-1.6.3.tar.bz2  
libksba-1.3.3.tar.bz2  
libassuan-2.2.0.tar.bz2  
pinentry-0.9.1.tar.bz2  
gpgme-1.5.3.tar.bz2  
gpa-0.9.7.tar.bz2  
dirmngr-1.1.0.tar.bz2

# Any alternatives to hashing?

- Protect the path
  - Lock it down, seal it up, etc.
- Detect tampering
  - Power loss, other physical changes
- All are very hard to do

# Encryption

- Convert an easily readable bit pattern into a bit pattern that looks very different
- Typically one that looks like random data
- Usually in a reversible way
  - So those you want to use the data can
  - Requires that not everyone can reverse it
- How to achieve that?

# Keyed encryption

- Use a secret to perform the conversion
- If you know the secret, reversing it is easy
- If you don't know the secret, reversing it is hard
  - Preferably impossible
- The secret is called the *key*
- Leading to an obvious question:
  - How can I keep a secret by using another secret?

# Symmetric and Asymmetric Encryption Systems

everyone knows my public key

encrypt with my private, and your public

- Symmetric systems use the same keys to encrypt and decrypt

decrypt with your private key and my public

  - Encrypt your data with key  $K$
  - Decrypt and get the data back with  $K$
- Asymmetric systems use different keys to encrypt and decrypt
  - Encrypt your data with  $K_E$
  - Decrypt and get the data back with  $K_D$
  - $K_E \neq K_D$ 
    - encrypt private, decrypt public to authenticate
    - authenticate my message - see who sent it
    - encrypt public, decrypt private to send

closed book closed notes  
no questions about 'do you know what this means...'  
apply knowledge to practical problems  
cumulative - write a couple of paragraphs - a page of text  
no electronic devices in the test

concentration in the class - both readings and lectures  
no need to memorize equations

# Example codes


- Symmetric
  - Data Encryption Standard (DES)
  - Advanced Encryption Standard (AES)
- Asymmetric
  - Diffie-Hellman
    - They just won the Turing Award for inventing asymmetric crypto
  - RSA algorithm
  - Elliptic curve algorithms

# Symmetric keys

- Also known as “shared secret”
  - Both sides share the same key
  - Both sides can encrypt or decrypt
- Generally faster than asymmetric crypto



This turns out to be really important!



# Using symmetric keys


- Assume the data you want to encrypt is  $P$
- The encryption algorithm is  $E$
- The decryption algorithm is  $D$
- And the symmetric key is  $K$
- $C = E(K, P)$
- $P = D(K, C)$
- Expanding,  $P = D(K, E(K, P))$
- You end up with what you started with
- And you used the same key twice



# Asymmetric keys

- *Public key cryptography*
  - Two keys: public and private
- To encrypt to a single recipient:
  - Anyone encrypts a message with your public key
  - Only you can decrypt with your private key
  - Only you can read it
- To identify a source:
  - You encrypt a message with your private key
  - Anyone can decrypt with your public key
  - Only you could have written it

ASYMMETRIC ENCRYPTION



# Symmetric vs. asymmetric

- Symmetric: same key is used on both ends
  - Anyone who has the key can create the message
  - Anyone who has the key can read the message
  - Info is private to those who share the key
  - Info was created by someone who knew the key
- Asymmetric: keys are used as pairs
  - Public key creates message only private key can decrypt
    - Confidentiality – only private key owner can read it
  - Private key creates message only public key can decrypt
    - Authenticity – only private key owner could create it
    - But anyone can check ownership
- Again, symmetric is much cheaper than asymmetric

# Using asymmetric keys

- Applying both keys yields the original message
  - $C = E(K_E, P)$
  - $P = D(K_D, C)$
- Or
  - $C = E(K_D, P)$
  - $P = D(K_E, C)$
- Unlike symmetric keys, the intermediates are different
  - $E(K_D, P) \neq E(K_E, P)$

# Digital signatures

- Rely on asymmetric keys
  - Signer encrypts using their private key
- Entire message?
  - That's too costly
    - Remember asymmetric being expensive?
  - Less costly to sign a hash
- Signature
  - A cryptographic hash signed with a private key
  - A.k.a. Message Authentication Code (MAC)

# Signatures and integrity

- Signature assures receiver of message integrity
  - Via the hash
- Contents haven't changed since the hash was computed
  - If they had, the hash wouldn't match
- Attacker can't just generate a new hash
  - Since it must be signed by the private key
  - Which he doesn't have
    - We hope . . .

# Signatures and authenticity

- Signature assures receiver of authenticity
  - Message was created by the apparent sender
- Sender's private key was used to sign the hash
  - Hash came from the party with *that* private key
  - Which can only be the apparent sender
    - We hope . . .

# Signatures and non-repudiation

- Signature prevents sender repudiation
  - Sender can't deny it sent that message
- Why not?
  - The decrypted signature matches the message hash
  - But it's a cryptographic hash
    - So it's not likely the message could be changed to match the signature
    - OR that a signature can be reused for a different message

# What do we have so far?

- Hash
  - Integrity, if you trust the hash
- Encryption
  - Privacy, given a key
- Signature
  - Authentication and integrity, given a key



# “given a key”

- Security’s three most feared words
  - Keys need to be shared in advance
    - Both sides have the symmetric key
    - Both sides have part of an asymmetric key
- The two challenges:
  - Endpoints need to know which key to use
    - Using the wrong key ruins everything
  - Endpoints need to get (and trust) the key
    - And if anyone else gets a symmetric key, you’re screwed

# Key management

- Pre-shared
- PKI
- Key exchange
- Keyless

# Pre-shared

- Both sides share the key in advance
  - Technically, this is usually assumed
  - Typically referred to as “out-of-band” distribution
    - Out-of-band = someone else’s job
    - I.e., “I’m not solving the hard part of the problem”
  - Useful for any keys
    - Shared secret (symmetric)
    - Public key (asymmetric)



© Colorpix.be

# PKI



- Public Key Infrastructure
  - Using a database to get a key
  - Same as most other network databases
    - Distributed vs. central
    - Flat vs. hierarchical
    - Structured vs. “hash tree” (destroys locality)
  - Infrastructure for public keys
    - Useful only for public part of asymmetric keys
    - Not a “public infrastructure for keys”

# PKI example

- PGP keys (e-mail)
  - Set of servers hold public keys
  - Users find keys explicitly
- X.509 keys
  - Key signed by a hierarchy
  - Many roots (built-in to browsers)
  - Can add others (self-signed, other-signed)

# PGP vs. X.509

- PGP
  - Web of trust
  - Users sign each other's keys
  - Key signing “parties”
  - Trust based on who YOU trust
- X.509
  - Hierarchy of trust
  - Roots sign keys
  - Companies charge to sign keys
  - Trust based on “anchors” (roots)



# Key Exchange

- Let's say you want to communicate
- But you don't share a key
  - Or you want to use a new key
  - Generally good not to use a single key too much
- Then you need to exchange a key between the communicating partners
- How?


# Key exchange using both symmetric and asymmetric crypto

- Common to use both in a single session
- Asymmetric cryptography essentially used to “bootstrap” symmetric crypto
- Use RSA (or another PK algorithm) to authenticate and establish a *session key*
- Use AES with that session key for the rest of the transmission



# Combining Symmetric and Asymmetric Crypto

Alice wants to share the key only  
with Bob



Unfortunately, it's more complex than this



Alice

$K_{EA}$

$K_{DA}$

$K_{EB}$

$K_S$

$C = E(K_S, K_{EB})$

$M = E(C, K_{DA})$

Only Bob can  
decrypt it

Only Alice could have created  
it

Take CS  
136 if  
you'd like  
to know  
why



Bob

$K_{EB}$


$K_{DB}$

$K_{EA}$

$M = D(C, K_{DB}) = D(M, K_{EA})$

# Diffie-Hellman key exchange

- Share a key starting from nothing
- Diffie-Hellman
  - Establish a shared secret over a public net
  - Each side has a secret (pick a random number)
  - Both sides share a common value
  - Relies on non-inverting mixing



# How to mix irreversibly?

- **Multiply**

- $A \times B = C$
- Assumes A and B are prime
- Factorization is hard

- **Exponentiation**

- $(a^x)^y \bmod p = (a^y)^x \bmod p$
- Discrete logarithms are hard

# What does DH do?

- DH establishes a shared secret
  - A symmetric key
- But symmetric keys don't establish identity
  - “Man in the middle” attack
  - Who do you share the secret with?
  - Solution: signed DH

# Signed DH

- Use public key cryptography
  - Sign the messages (encrypt with private key)
  - Prevents MITM attack
- But if we have public key, why use not use that?
  - Remember how important it was that asymmetric crypto was expensive?
- Why not what we did a few slides ago?
  - Requires fewer of those expensive PK operations

# Keyless

- The I in PKI is a pain
  - The hardest part of PKI is the key database
  - Everyone has to have a key
  - Everyone has to find the other party's key
- Solutions
  - Make the “I” easier (automate, etc.)
  - Avoid the “PK”

# Keyless keying

- What if we use DH without signatures?
  - I.e., original DH
  - Share a secret – but with whom?
- What if we don't care?
  - WHO isn't know
  - But the rest of the exchange is protected
- “Better than nothing security” (BTNS)
  - Protects against others interfering

# BTNS protection

- The connection is secure!
  - But to whom?
    - Who cares!
    - Once you start a conversation, you can't be interrupted
    - Maybe you can somehow verify identity later?



# Are we done?

- What have we protected?
  - Integrity
  - Privacy
  - Origin
- What have we not protected?
  - Resources!

# Resource protection

- Endpoint
- Forwarding

# Endpoint resource protection

- Resources to protect
  - Buffers *memory*
  - Processing / CPU *processing*
  - Content – the FSMs
- Ways to protect the endpoint
  - Shed load
  - Verify before acting

hard to figure out everything you should drop

drop traffic - drop all the junk; don't drop a single packet of what you don't want

# Typical endpoint protections

- Rate limiting

firewall - sits at a place where it checks packets  
get rid of traffic that you don't want your server to be handling  
examine the incoming traffic - see which one you like

- Limit investment in new connections
- Toss out when beyond a limit
- Protect against SYN flooding attacks

- Firewalls, port blocking

- Fixed: drop all packets to a particular port and/or in a particular direction
- Conditional: drop a port until you know better

SYN flood - TCP SYN's are sent, but never finishes the connection  
limit your resources, you don't want incoming traffic to set up resource for connection

# Conditional port blocking

- NATs
  - Network and port address translator
  - Private and public side
  - Fixed: public side -> private side
  - Conditional: private side -> public side
- Conditional example
  - Allow incoming only if outgoing
    - Wait for DNS UDP out, allow response back in
    - Wait for TCP SYN (open), allow replies until FIN (close)

encrypt internal packet - wrap with another IP header - this is from IP address in starbucks  
gets delivered to UCLA tunneling layer - decrypts and get IP packet

# Variable load shedding

load shedding - too much work, time to drop something

- Port blocking: drop based on partial work
  - Examine addresses, ports, some content
  - Drop before investing more work
- Cipher/code/sign: drop on separate work
  - Validate (decrypt, authenticate) based on an algorithm that is separate from the FSM of the protocol
  - Drop before performing separate work
- Both attempt to separate security from FSM
  - Checking is distinct from acting on the message

exception to encrypt header  
link - important  
tunneling - have protocol  
hide the protocol we want to  
run, throw another layer  
on top of things

tunnel IP on top of HTTP  
encapsulate IP on top of  
HTTP

VPN - tunneling

don't encrypt the IP header on packets -  
tunneling always has some sort of cost

runnel any protocol on top of another protocl - tunnel UDP on top of TCP

# Forwarding resource protection

- Routers have two distinct roles
  - Relaying messages
  - As endpoints of routing and control protocols
- Two kinds of protection
  - Endpoint-like
  - Forwarding focused

# Router endpoint protection

- Similar to other endpoints
  - Block ports
  - Limit rate
  - Validate content
- But a little harder sometimes
  - Bellman-Ford relays content indirectly
    - How can you protect the FSM?
    - Do you attach signatures for all the path components?
    - How does this affect scalability?



# Forwarding protection

- Why would I refuse to forward a packet?
  - Others similar packets are causing a problem here
    - Overloading of security processing
    - Overloading my buffers
  - Other similar packets are causing a problem elsewhere (most commonly downstream)
    - The other end of the link has no room for it
    - The other end of the link says these are a problem
  - It never should have come to me
    - “Reverse path” checks

DDoS attacks - call sys admin  
analyze your traffic

# Where do we put security?

- Everywhere you want to protect resources
  - Everywhere in the DAG
  - Sometimes at every hop, sometimes on ends
- In layers
  - Always better to shed load earlier in traversal
  - Don't always have enough info until later

# Summary

- Security happens everywhere
  - At every layer
  - As early as possible, but no earlier
- Security uses several tools
  - Most based on math
  - Not all based on pure math – lots of alchemy
- Security assumes shared secrets
  - Like naming, it can't start from nothing